

knn

April 16, 2022

```
[1]: # This mounts your Google Drive to the Colab VM.
from google.colab import drive
drive.mount('/content/drive')

# TODO: Enter the foldername in your Drive where you have saved the unzipped
# assignment folder, e.g. 'cs231n/assignments/assignment1/'
FOLDERNAME = 'Colab_Notebooks/cs231n/assignments/assignment1/'
assert FOLDERNAME is not None, "[!] Enter the foldername."

# Now that we've mounted your Drive, this ensures that
# the Python interpreter of the Colab VM can load
# python files from within it.
import sys
sys.path.append('/content/drive/My Drive/{}'.format(FOLDERNAME))

# This downloads the CIFAR-10 dataset to your Drive
# if it doesn't already exist.
%cd /content/drive/My\ Drive/$FOLDERNAME/cs231n/datasets/
!bash get_datasets.sh
%cd /content/drive/My\ Drive/$FOLDERNAME
```

```
Mounted at /content/drive
/content/drive/My
Drive/Colab_Notebooks/cs231n/assignments/assignment1/cs231n/datasets
/content/drive/My Drive/Colab_Notebooks/cs231n/assignments/assignment1
```

1 k-Nearest Neighbor (kNN) exercise

Complete and hand in this completed worksheet (including its outputs and any supporting code outside of the worksheet) with your assignment submission. For more details see the [assignments page](#) on the course website.

The kNN classifier consists of two stages:

- During training, the classifier takes the training data and simply remembers it
- During testing, kNN classifies every test image by comparing to all training images and transferring the labels of the k most similar training examples

- The value of k is cross-validated

In this exercise you will implement these steps and understand the basic Image Classification pipeline, cross-validation, and gain proficiency in writing efficient, vectorized code.

```
[2]: # Run some setup code for this notebook.

import random
import numpy as np
from cs231n.data_utils import load_CIFAR10
import matplotlib.pyplot as plt

# This is a bit of magic to make matplotlib figures appear inline in the
↪notebook
# rather than in a new window.
%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# Some more magic so that the notebook will reload external python modules;
# see http://stackoverflow.com/questions/1907993/
↪autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

[3]: # Load the raw CIFAR-10 data.
cifar10_dir = 'cs231n/datasets/cifar-10-batches-py'

# Cleaning up variables to prevent loading data multiple times (which may cause
↪memory issue)
try:
    del X_train, y_train
    del X_test, y_test
    print('Clear previously loaded data.')
except:
    pass

X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)

# As a sanity check, we print out the size of the training and test data.
print('Training data shape: ', X_train.shape)
print('Training labels shape: ', y_train.shape)
print('Test data shape: ', X_test.shape)
print('Test labels shape: ', y_test.shape)
```

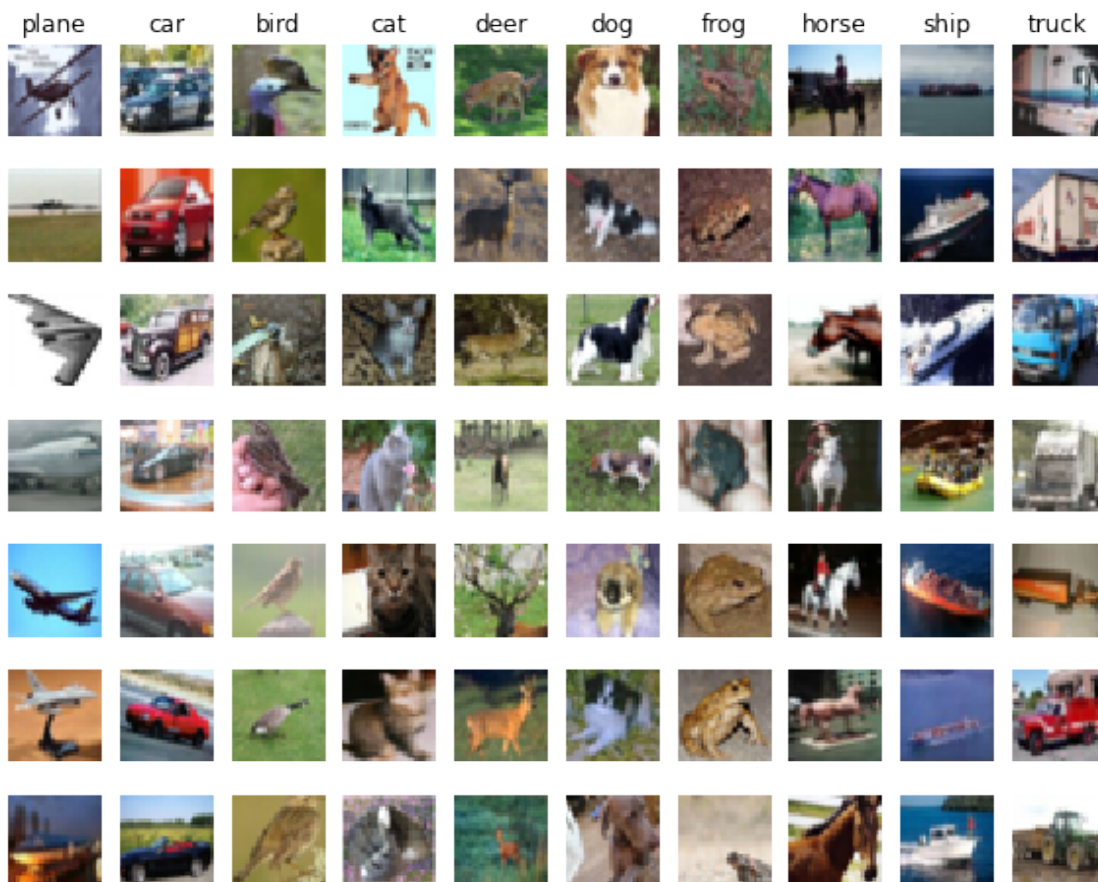
Training data shape: (50000, 32, 32, 3)

Training labels shape: (50000,)

Test data shape: (10000, 32, 32, 3)

Test labels shape: (10000,)

```
[4]: # Visualize some examples from the dataset.
# We show a few examples of training images from each class.
classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
num_classes = len(classes)
samples_per_class = 7
for y, cls in enumerate(classes):
    idxs = np.flatnonzero(y_train == y)
    idxs = np.random.choice(idxs, samples_per_class, replace=False)
    for i, idx in enumerate(idxs):
        plt_idx = i * num_classes + y + 1
        plt.subplot(samples_per_class, num_classes, plt_idx)
        plt.imshow(X_train[idx].astype('uint8'))
        plt.axis('off')
    if i == 0:
        plt.title(cls)
plt.show()
```



```
[5]: # Subsample the data for more efficient code execution in this exercise
num_training = 5000
mask = list(range(num_training))
X_train = X_train[mask]
y_train = y_train[mask]

num_test = 500
mask = list(range(num_test))
X_test = X_test[mask]
y_test = y_test[mask]

# Reshape the image data into rows
X_train = np.reshape(X_train, (X_train.shape[0], -1))
X_test = np.reshape(X_test, (X_test.shape[0], -1))
print(X_train.shape, X_test.shape)
```

(5000, 3072) (500, 3072)

```
[6]: from cs231n.classifiers import KNearestNeighbor

# Create a kNN classifier instance.
# Remember that training a kNN classifier is a noop:
# the Classifier simply remembers the data and does no further processing
classifier = KNearestNeighbor()
classifier.train(X_train, y_train)
```

We would now like to classify the test data with the kNN classifier. Recall that we can break down this process into two steps:

1. First we must compute the distances between all test examples and all train examples.
2. Given these distances, for each test example we find the k nearest examples and have them vote for the label

Lets begin with computing the distance matrix between all training and test examples. For example, if there are **N_{tr}** training examples and **N_{te}** test examples, this stage should result in a **N_{te} x N_{tr}** matrix where each element (i,j) is the distance between the i-th test and j-th train example.

Note: For the three distance computations that we require you to implement in this notebook, you may not use the `np.linalg.norm()` function that numpy provides.

First, open `cs231n/classifiers/k_nearest_neighbor.py` and implement the function `compute_distances_two_loops` that uses a (very inefficient) double loop over all pairs of (test, train) examples and computes the distance matrix one element at a time.


```
[7]: # Open cs231n/classifiers/k_nearest_neighbor.py and implement
# compute_distances_two_loops.

# Test your implementation:
```

```
dists = classifier.compute_distances_two_loops(X_test)
print(dists.shape)
```

(500, 5000)

```
[8]: # We can visualize the distance matrix: each row is a single test example and
# its distances to training examples
plt.imshow(dists, interpolation='none')
plt.show()
```



Inline Question 1

Notice the structured patterns in the distance matrix, where some rows or columns are visibly brighter. (Note that with the default color scheme black indicates low distances while white indicates high distances.)

- What in the data is the cause behind the distinctly bright rows?
- What causes the columns?

Your Answer : \

Q 1: Bright rows correspond to the high distance between a given test example and most of the training examples because of background color, illumination and others. For example: image of a ship (with majority of blue color) has consistently higher distance for all classes except ship and aeroplane. Some image with white background has higher distances with training images with non white background. \

Q 2: Similar reasoning can be found for the columns with white color, where a training example of ship/aeroplane has consistently higher distance with all classes except ship and aeroplane.

```
[9]: # Now implement the function predict_labels and run the code below:
# We use k = 1 (which is Nearest Neighbor).
y_test_pred = classifier.predict_labels(dists, k=1)

# Compute and print the fraction of correctly predicted examples
num_correct = np.sum(y_test_pred == y_test)
accuracy = float(num_correct) / num_test
print('Got %d / %d correct => accuracy: %f' % (num_correct, num_test, accuracy))
```

Got 137 / 500 correct => accuracy: 0.274000

You should expect to see approximately 27% accuracy. Now lets try out a larger k, say k = 5:

```
[10]: y_test_pred = classifier.predict_labels(dists, k=5)
num_correct = np.sum(y_test_pred == y_test)
accuracy = float(num_correct) / num_test
print('Got %d / %d correct => accuracy: %f' % (num_correct, num_test, accuracy))
```

Got 139 / 500 correct => accuracy: 0.278000

You should expect to see a slightly better performance than with $k = 1$.

Inline Question 2

We can also use other distance metrics such as L1 distance. For pixel values $p_{ij}^{(k)}$ at location (i, j) of some image I_k ,

the mean μ across all pixels over all images is

$$\mu = \frac{1}{nhw} \sum_{k=1}^n \sum_{i=1}^h \sum_{j=1}^w p_{ij}^{(k)}$$

And the pixel-wise mean μ_{ij} across all images is

$$\mu_{ij} = \frac{1}{n} \sum_{k=1}^n p_{ij}^{(k)}.$$

The general standard deviation σ and pixel-wise standard deviation σ_{ij} is defined similarly.

Which of the following preprocessing steps will not change the performance of a Nearest Neighbor classifier that uses L1 distance? Select all that apply. 1. Subtracting the mean μ ($\tilde{p}_{ij}^{(k)} = p_{ij}^{(k)} - \mu$.) 2. Subtracting the per pixel mean μ_{ij} ($\tilde{p}_{ij}^{(k)} = p_{ij}^{(k)} - \mu_{ij}$.) 3. Subtracting the mean μ and dividing by the standard deviation σ . 4. Subtracting the pixel-wise mean μ_{ij} and dividing by the pixel-wise standard deviation σ_{ij} . 5. Rotating the coordinate axes of the data.

Your Answer : \ Step 1, 2, 3 will not affect the result of k-NN \

Your Explanation : \ Reason for Step 1: Subtracting a constant value (mean) from all pixel will not change the L1 distance as it is similar to shift of origin which doesn't affect the distance between two points. \ Reason for Step 2: Subtracting the pixel wise mean across all images will ensure L1 distance to be constant as before. Reason being same: shift of origin for each pixel with respect to different value will cancel out when it is subtracted from test image's pixel value. \ Reason for Step 3: L1 distance remains unchanged because entire dataset will be normalized w.r.t. mean and standard deviation. It is data normalization step for making data centered around 0 and with standard deviation 1. Hence L1 distance would remain unchanged. \

```
[11]: # Now lets speed up distance matrix computation by using partial vectorization
# with one loop. Implement the function compute_distances_one_loop and run the
# code below:
dists_one = classifier.compute_distances_one_loop(X_test)

# To ensure that our vectorized implementation is correct, we make sure that it
# agrees with the naive implementation. There are many ways to decide whether
# two matrices are similar; one of the simplest is the Frobenius norm. In case
```

```

# you haven't seen it before, the Frobenius norm of two matrices is the square
# root of the squared sum of differences of all elements; in other words,
→ reshape
# the matrices into vectors and compute the Euclidean distance between them.
difference = np.linalg.norm(dists - dists_one, ord='fro')
print('One loop difference was: %f' % (difference, ))
if difference < 0.001:
    print('Good! The distance matrices are the same')
else:
    print('Uh-oh! The distance matrices are different')

```

One loop difference was: 0.000000
Good! The distance matrices are the same

```

[12]: # Now implement the fully vectorized version inside compute_distances_no_loops
# and run the code
dists_two = classifier.compute_distances_no_loops(X_test)

# check that the distance matrix agrees with the one we computed before:
difference = np.linalg.norm(dists - dists_two, ord='fro')
print('No loop difference was: %f' % (difference, ))
if difference < 0.001:
    print('Good! The distance matrices are the same')
else:
    print('Uh-oh! The distance matrices are different')

```

No loop difference was: 0.000000
Good! The distance matrices are the same

```

[13]: # Let's compare how fast the implementations are
def time_function(f, *args):
    """
    Call a function f with args and return the time (in seconds) that it took
    → to execute.
    """
    import time
    tic = time.time()
    f(*args)
    toc = time.time()
    return toc - tic

two_loop_time = time_function(classifier.compute_distances_two_loops, X_test)
print('Two loop version took %f seconds' % two_loop_time)

one_loop_time = time_function(classifier.compute_distances_one_loop, X_test)
print('One loop version took %f seconds' % one_loop_time)

```

```

no_loop_time = time_function(classifier.compute_distances_no_loops, X_test)
print('No loop version took %f seconds' % no_loop_time)

# You should see significantly faster performance with the fully vectorized
  ↳ implementation!

# NOTE: depending on what machine you're using,
# you might not see a speedup when you go from two loops to one loop,
# and might even see a slow-down.

```

Two loop version took 43.269078 seconds
 One loop version took 38.155360 seconds
 No loop version took 0.620053 seconds

1.0.1 Cross-validation

We have implemented the k-Nearest Neighbor classifier but we set the value $k = 5$ arbitrarily. We will now determine the best value of this hyperparameter with cross-validation.

```

[14]: num_folds = 5
      k_choices = [1, 3, 5, 8, 10, 12, 15, 20, 50, 100]

      X_train_folds = []
      y_train_folds = []
      #####
      # TODO:
      # Split up the training data into folds. After splitting, X_train_folds and
      # y_train_folds should each be lists of length num_folds, where
      # y_train_folds[i] is the label vector for the points in X_train_folds[i].
      # Hint: Look up the numpy array_split function.
      #####
      # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
      X_train_folds = np.array_split(X_train, num_folds)
      y_train_folds = np.array_split(y_train, num_folds)
      #print(X_train_folds[0].shape, y_train_folds[0].shape)
      # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

      # A dictionary holding the accuracies for different values of k that we find
      # when running cross-validation. After running cross-validation,
      # k_to_accuracies[k] should be a list of length num_folds giving the different
      # accuracy values that we found when using that value of k.
      k_to_accuracies = {}

      #####
      # TODO:
      # Perform k-fold cross validation to find the best value of k. For each
      # possible value of k, run the k-nearest-neighbor algorithm num_folds times,

```



```

# where in each case you use all but one of the folds as training data and the #
# last fold as a validation set. Store the accuracies for all fold and all      #
# values of k in the k_to_accuracies dictionary.                               #
#####
# *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

for k in k_choices:
    accuracy = []
    for i in range(num_folds):
        #print(k , i, X_train_folds[i].shape)

        # creating the X_val and y_val data
        X_val = X_train_folds[i]
        y_val = y_train_folds[i]

        # form the X_train, y_train using all but i_th X_train_folds[i] using
        ↪ concatenating
        X_train_list = X_train_folds[:i] + X_train_folds[i+1:]
        y_train_list = y_train_folds[:i] + y_train_folds[i+1:]

        # train the classifier
        X_train_cross = np.concatenate(X_train_list)
        y_train_cross = np.concatenate(y_train_list)
        classifier.train(X_train_cross, y_train_cross)

        dists = classifier.compute_distances_no_loops(X_val)
        y_val_pred = classifier.predict_labels(dists, k=k)
        num_correct = np.sum(y_val_pred == y_val)

        # compute accuracy and append it to a list
        accuracy.append(float(num_correct) / y_val.shape[0])
    k_to_accuracies[k] = accuracy

# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

# Print out the computed accuracies
for k in sorted(k_to_accuracies):
    for accuracy in k_to_accuracies[k]:
        print('k = %d, accuracy = %f' % (k, accuracy))

```

```


k = 1, accuracy = 0.263000
k = 1, accuracy = 0.257000
k = 1, accuracy = 0.264000
k = 1, accuracy = 0.278000
k = 1, accuracy = 0.266000

```

k = 3, accuracy = 0.239000
k = 3, accuracy = 0.249000
k = 3, accuracy = 0.240000
k = 3, accuracy = 0.266000
k = 3, accuracy = 0.254000
k = 5, accuracy = 0.248000
k = 5, accuracy = 0.266000
k = 5, accuracy = 0.280000
k = 5, accuracy = 0.292000
k = 5, accuracy = 0.280000
k = 8, accuracy = 0.262000
k = 8, accuracy = 0.282000
k = 8, accuracy = 0.273000
k = 8, accuracy = 0.290000
k = 8, accuracy = 0.273000
k = 10, accuracy = 0.265000
k = 10, accuracy = 0.296000
k = 10, accuracy = 0.276000
k = 10, accuracy = 0.284000
k = 10, accuracy = 0.280000
k = 12, accuracy = 0.260000
k = 12, accuracy = 0.295000
k = 12, accuracy = 0.279000
k = 12, accuracy = 0.283000
k = 12, accuracy = 0.280000
k = 15, accuracy = 0.252000
k = 15, accuracy = 0.289000
k = 15, accuracy = 0.278000
k = 15, accuracy = 0.282000
k = 15, accuracy = 0.274000
k = 20, accuracy = 0.270000
k = 20, accuracy = 0.279000
k = 20, accuracy = 0.279000
k = 20, accuracy = 0.282000
k = 20, accuracy = 0.285000
k = 50, accuracy = 0.271000
k = 50, accuracy = 0.288000
k = 50, accuracy = 0.278000
k = 50, accuracy = 0.269000
k = 50, accuracy = 0.266000
k = 100, accuracy = 0.256000
k = 100, accuracy = 0.270000
k = 100, accuracy = 0.263000
k = 100, accuracy = 0.256000
k = 100, accuracy = 0.263000

```
[15]: # plot the raw observations
for k in k_choices:
    accuracies = k_to_accuracies[k]
    plt.scatter([k] * len(accuracies), accuracies)

# plot the trend line with error bars that correspond to standard deviation
accuracies_mean = np.array([np.mean(v) for k,v in sorted(k_to_accuracies.
    ↳items())])
accuracies_std = np.array([np.std(v) for k,v in sorted(k_to_accuracies.
    ↳items())])
plt.errorbar(k_choices, accuracies_mean, yerr=accuracies_std)
plt.title('Cross-validation on k')
plt.xlabel('k')
plt.ylabel('Cross-validation accuracy')
plt.show()
```



```
[16]: # Based on the cross-validation results above, choose the best value for k,
# retrain the classifier using all the training data, and test it on the test
# data. You should be able to get above 28% accuracy on the test data.
```

```

best_k = k_choices[accuracies_mean.argmax()]

classifier = KNearestNeighbor()
classifier.train(X_train, y_train)
y_test_pred = classifier.predict(X_test, k=best_k)

# Compute and display the accuracy
num_correct = np.sum(y_test_pred == y_test)
accuracy = float(num_correct) / num_test
print('Got %d / %d correct => accuracy: %f' % (num_correct, num_test, accuracy))

```

Got 141 / 500 correct => accuracy: 0.282000

Inline Question 3

Which of the following statements about k -Nearest Neighbor (k -NN) are true in a classification setting, and for all k ? Select all that apply. 1. The decision boundary of the k -NN classifier is linear. 2. The training error of a 1-NN will always be lower than or equal to that of 5-NN. 3. The test error of a 1-NN will always be lower than that of a 5-NN. 4. The time needed to classify a test example with the k -NN classifier grows with the size of the training set. 5. None of the above.

Your Answer : \ Option 1: False \ Option 2: True \ Option 3: False \ Option 4: True \

Your Explanation : \ Option 1: As the value of k increases, generalizability of the model increases and leads to smoother decision boundary and hence decision boundary tends to become non-linear. \ Option 2: During training, input consists of both x, y ; so during training 1-NN will have 100% training accuracy (as the 1-NN remembers all the data points) but will have lower test accuracy than or equal to 5-NN. \ Option 3: During test time, input data only contains x , so 1-NN considers only 1-NN as opposed to 5-NN which has more smooth decision boundary. Hence 1-NN test error can be less than or equal to 5-NN, but NOT always. \ Option 4: Inference using k -NN involves comparing L1/L2 distance for each test (out of total M test images) to each train images (out of N total test image), hence inference time increases with $O(N)$ as size of training set increases.

SVM

April 16, 2022

```
[19]: # This mounts your Google Drive to the Colab VM.
from google.colab import drive
drive.mount('/content/drive')

# TODO: Enter the foldername in your Drive where you have saved the unzipped
# assignment folder, e.g. 'cs231n/assignments/assignment1/'
FOLDERNAME = 'Colab_Notebooks/cs231n/assignments/assignment1/'
assert FOLDERNAME is not None, "[!] Enter the foldername."

# Now that we've mounted your Drive, this ensures that
# the Python interpreter of the Colab VM can load
# python files from within it.
import sys
sys.path.append('/content/drive/My Drive/{}'.format(FOLDERNAME))

# This downloads the CIFAR-10 dataset to your Drive
# if it doesn't already exist.
%cd /content/drive/My\ Drive/$FOLDERNAME/cs231n/datasets/
!bash get_datasets.sh
%cd /content/drive/My\ Drive/$FOLDERNAME
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call `drive.mount("/content/drive", force_remount=True)`.

/content/drive/My

Drive/Colab_Notebooks/cs231n/assignments/assignment1/cs231n/datasets

/content/drive/My Drive/Colab_Notebooks/cs231n/assignments/assignment1

1 Multiclass Support Vector Machine exercise

Complete and hand in this completed worksheet (including its outputs and any supporting code outside of the worksheet) with your assignment submission. For more details see the [assignments page](#) on the course website.

In this exercise you will:

- implement a fully-vectorized **loss function** for the SVM
- implement the fully-vectorized expression for its **analytic gradient**

- **check your implementation** using numerical gradient
- use a validation set to **tune the learning rate and regularization** strength
- **optimize** the loss function with **SGD**
- **visualize** the final learned weights

```
[20]: # Run some setup code for this notebook.
import random
import numpy as np
from cs231n.data_utils import load_CIFAR10
import matplotlib.pyplot as plt

# This is a bit of magic to make matplotlib figures appear inline in the
# notebook rather than in a new window.
%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# Some more magic so that the notebook will reload external python modules;
# see http://stackoverflow.com/questions/1907993/
# ↪ autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2
```

The autoreload extension is already loaded. To reload it, use:

```
%reload_ext autoreload
```

1.1 CIFAR-10 Data Loading and Preprocessing

```
[21]: # Load the raw CIFAR-10 data.
cifar10_dir = 'cs231n/datasets/cifar-10-batches-py'

# Cleaning up variables to prevent loading data multiple times (which may cause
# ↪ memory issue)
try:
    del X_train, y_train
    del X_test, y_test
    print('Clear previously loaded data.')
except:
    pass

X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)

# As a sanity check, we print out the size of the training and test data.
print('Training data shape: ', X_train.shape)
print('Training labels shape: ', y_train.shape)
```

```
print('Test data shape: ', X_test.shape)
print('Test labels shape: ', y_test.shape)
```

Clear previously loaded data.

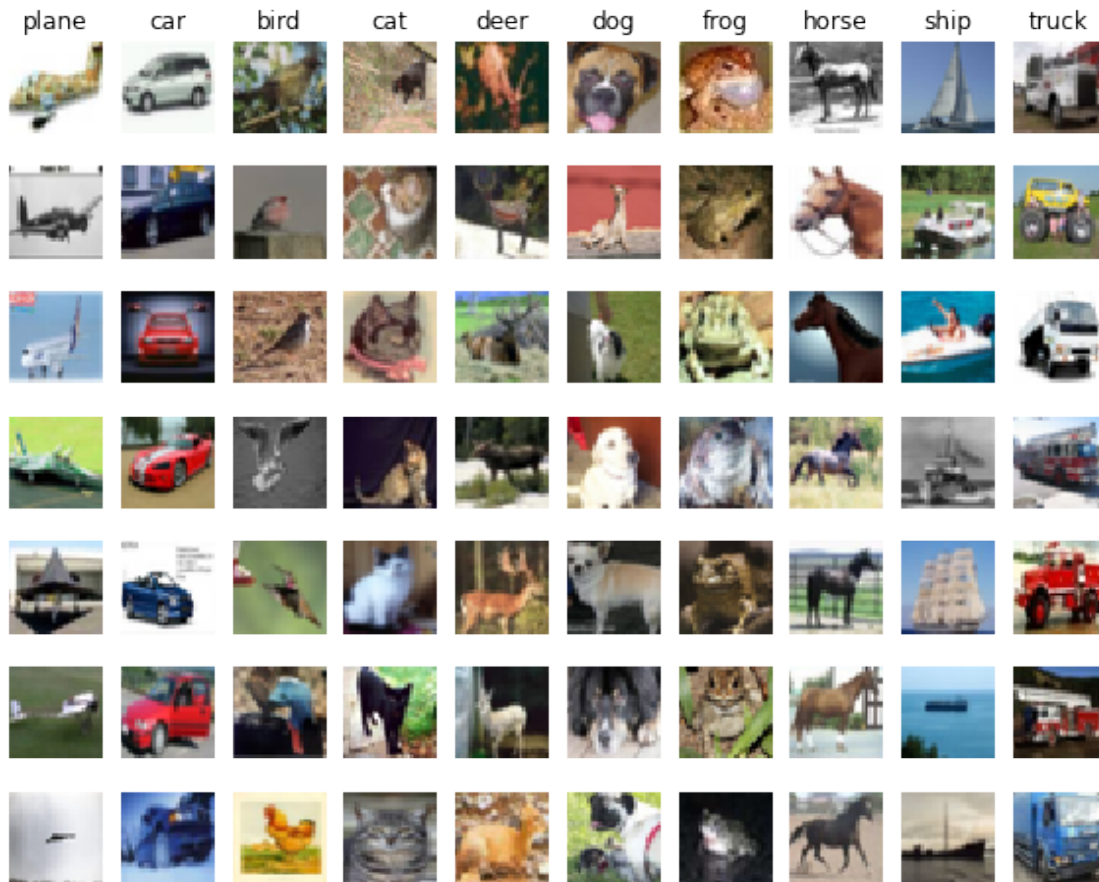
Training data shape: (50000, 32, 32, 3)

Training labels shape: (50000,)

Test data shape: (10000, 32, 32, 3)

Test labels shape: (10000,)

```
[22]: # Visualize some examples from the dataset.
# We show a few examples of training images from each class.
classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', '
↳ 'ship', 'truck']
num_classes = len(classes)
samples_per_class = 7
for y, cls in enumerate(classes):
    idxs = np.flatnonzero(y_train == y)
    idxs = np.random.choice(idxs, samples_per_class, replace=False)
    for i, idx in enumerate(idxs):
        plt_idx = i * num_classes + y + 1
        plt.subplot(samples_per_class, num_classes, plt_idx)
        plt.imshow(X_train[idx].astype('uint8'))
        plt.axis('off')
        if i == 0:
            plt.title(cls)
plt.show()
```



```
[23]: # Split the data into train, val, and test sets. In addition we will
# create a small development set as a subset of the training data;
# we can use this for development so our code runs faster.
num_training = 49000
num_validation = 1000
num_test = 1000
num_dev = 500

# Our validation set will be num_validation points from the original
# training set.
mask = range(num_training, num_training + num_validation)
X_val = X_train[mask]
y_val = y_train[mask]

# Our training set will be the first num_train points from the original
# training set.
mask = range(num_training)
X_train = X_train[mask]
y_train = y_train[mask]
```



```

# We will also make a development set, which is a small subset of
# the training set.
mask = np.random.choice(num_training, num_dev, replace=False)
X_dev = X_train[mask]
y_dev = y_train[mask]

# We use the first num_test points of the original test set as our
# test set.
mask = range(num_test)
X_test = X_test[mask]
y_test = y_test[mask]

print('Train data shape: ', X_train.shape)
print('Train labels shape: ', y_train.shape)
print('Validation data shape: ', X_val.shape)
print('Validation labels shape: ', y_val.shape)
print('Test data shape: ', X_test.shape)
print('Test labels shape: ', y_test.shape)

```

```

Train data shape: (49000, 32, 32, 3)
Train labels shape: (49000,)
Validation data shape: (1000, 32, 32, 3)
Validation labels shape: (1000,)
Test data shape: (1000, 32, 32, 3)
Test labels shape: (1000,)

```

```

[24]: # Preprocessing: reshape the image data into rows
X_train = np.reshape(X_train, (X_train.shape[0], -1))
X_val = np.reshape(X_val, (X_val.shape[0], -1))
X_test = np.reshape(X_test, (X_test.shape[0], -1))
X_dev = np.reshape(X_dev, (X_dev.shape[0], -1))

# As a sanity check, print out the shapes of the data
print('Training data shape: ', X_train.shape)
print('Validation data shape: ', X_val.shape)
print('Test data shape: ', X_test.shape)
print('dev data shape: ', X_dev.shape)

```

```

Training data shape: (49000, 3072)
Validation data shape: (1000, 3072)
Test data shape: (1000, 3072)
dev data shape: (500, 3072)

```

```

[25]: # Preprocessing: subtract the mean image
# first: compute the image mean based on the training data
mean_image = np.mean(X_train, axis=0)

```

```

print(mean_image[:10]) # print a few of the elements
plt.figure(figsize=(4,4))
plt.imshow(mean_image.reshape((32,32,3)).astype('uint8')) # visualize the mean_
    ↳image
plt.show()

# second: subtract the mean image from train and test data
X_train -= mean_image
X_val -= mean_image
X_test -= mean_image
X_dev -= mean_image

# third: append the bias dimension of ones (i.e. bias trick) so that our SVM
# only has to worry about optimizing a single weight matrix W.
X_train = np.hstack([X_train, np.ones((X_train.shape[0], 1))])
X_val = np.hstack([X_val, np.ones((X_val.shape[0], 1))])
X_test = np.hstack([X_test, np.ones((X_test.shape[0], 1))])
X_dev = np.hstack([X_dev, np.ones((X_dev.shape[0], 1))])


print(X_train.shape, X_val.shape, X_test.shape, X_dev.shape)

```

```

[130.64189796 135.98173469 132.47391837 130.05569388 135.34804082
 131.75402041 130.96055102 136.14328571 132.47636735 131.48467347]

```



```

(49000, 3073) (1000, 3073) (1000, 3073) (500, 3073)

```

1.2 SVM Classifier

Your code for this section will all be written inside `cs231n/classifiers/linear_svm.py`.

As you can see, we have prefilled the function `svm_loss_naive` which uses for loops to evaluate the multiclass SVM loss function.

```
[38]: # Evaluate the naive implementation of the loss we provided for you:
from cs231n.classifiers.linear_svm import svm_loss_naive
import time

# generate a random SVM weight matrix of small numbers
W = np.random.randn(3073, 10) * 0.0001

loss, grad = svm_loss_naive(W, X_dev, y_dev, 0.000005)
print('loss: %f' % (loss, ))
#print(grad)
```

loss: 9.808086

The `grad` returned from the function above is right now all zero. Derive and implement the gradient for the SVM cost function and implement it inline inside the function `svm_loss_naive`. You will find it helpful to interleave your new code inside the existing function.

To check that you have correctly implemented the gradient, you can numerically estimate the gradient of the loss function and compare the numeric estimate to the gradient that you computed. We have provided code that does this for you:

```
[27]: # Once you've implemented the gradient, recompute it with the code below
# and gradient check it with the function we provided for you

# Compute the loss and its gradient at W.
loss, grad = svm_loss_naive(W, X_dev, y_dev, 0.0)

# Numerically compute the gradient along several randomly chosen dimensions, and
# compare them with your analytically computed gradient. The numbers should
→match
# almost exactly along all dimensions.
from cs231n.gradient_check import grad_check_sparse
f = lambda w: svm_loss_naive(w, X_dev, y_dev, 0.0)[0]
grad_numerical = grad_check_sparse(f, W, grad)

# do the gradient check once again with regularization turned on
# you didn't forget the regularization gradient did you?
loss, grad = svm_loss_naive(W, X_dev, y_dev, 5e1)
f = lambda w: svm_loss_naive(w, X_dev, y_dev, 5e1)[0]
grad_numerical = grad_check_sparse(f, W, grad)
```

numerical: 26.796770 analytic: 26.796770, relative error: 1.378897e-12
numerical: 0.060142 analytic: 0.060142, relative error: 3.513350e-10

```

numerical: 13.087564 analytic: 13.087564, relative error: 3.081379e-12
numerical: 10.847550 analytic: 10.847550, relative error: 6.129406e-13
numerical: -5.447078 analytic: -5.447078, relative error: 1.590399e-11
numerical: -1.451508 analytic: -1.451508, relative error: 3.697333e-10
numerical: -21.175017 analytic: -21.175017, relative error: 8.455955e-12
numerical: 51.663508 analytic: 51.661455, relative error: 1.986710e-05
numerical: 27.133310 analytic: 27.133310, relative error: 1.116677e-11
numerical: 13.086252 analytic: 13.086252, relative error: 1.706415e-12
numerical: -8.790538 analytic: -8.790538, relative error: 1.956692e-11
numerical: 11.631147 analytic: 11.631147, relative error: 2.509868e-11
numerical: -22.330283 analytic: -22.270105, relative error: 1.349282e-03
numerical: -2.889024 analytic: -2.889024, relative error: 1.406052e-10
numerical: 8.249580 analytic: 8.249580, relative error: 1.892627e-11
numerical: 5.123546 analytic: 5.123546, relative error: 6.322071e-12
numerical: -7.535639 analytic: -7.535639, relative error: 4.325528e-11
numerical: 24.453916 analytic: 24.453916, relative error: 1.157098e-12
numerical: 8.200473 analytic: 8.200473, relative error: 1.740989e-11
numerical: 6.076467 analytic: 6.076467, relative error: 1.442843e-11

```

Inline Question 1

It is possible that once in a while a dimension in the gradcheck will not match exactly. What could such a discrepancy be caused by? Is it a reason for concern? What is a simple example in one dimension where a gradient check could fail? How would change the margin affect of the frequency of this happening? *Hint: the SVM loss function is not strictly speaking differentiable*

Your Answer : \ Because the max is non-differentiable, it has kinks in the loss function, since the gradient in these position is not defined. That is why making the delta bigger allow us to have less points with these kinks. A possible solution could be to use the subgradient. \ Also the numerical solution, is not exact, since it depends on h. if delta is 10e10 we are starting seen issues. But in general increasing delta will lower the frequency of kinks, we will have less negative values that could cause kinks when equal to 0.

```

[28]: # Next implement the function svm_loss_vectorized; for now only compute the
      ↪ loss;
      # we will implement the gradient in a moment.
      # https://mlxai.github.io/2017/01/06/
      ↪ vectorized-implementation-of-svm-loss-and-gradient-update.html
      tic = time.time()
      loss_naive, grad_naive = svm_loss_naive(W, X_dev, y_dev, 0.000005)
      toc = time.time()
      print('Naive loss: %e computed in %fs' % (loss_naive, toc - tic))

      from cs231n.classifiers.linear_svm import svm_loss_vectorized
      tic = time.time()
      loss_vectorized, _ = svm_loss_vectorized(W, X_dev, y_dev, 0.000005)
      toc = time.time()
      print('Vectorized loss: %e computed in %fs' % (loss_vectorized, toc - tic))

```

```
# The losses should match but your vectorized implementation should be much
→faster.
print('difference: %f' % (loss_naive - loss_vectorized))
```

Naive loss: 8.987904e+00 computed in 0.100972s
 Vectorized loss: 8.987904e+00 computed in 0.013593s
 difference: -0.000000

[29]: *# Complete the implementation of svm_loss_vectorized, and compute the gradient*
of the loss function in a vectorized way.

```
# The naive implementation and the vectorized implementation should match, but
# the vectorized version should still be much faster.
tic = time.time()
_, grad_naive = svm_loss_naive(W, X_dev, y_dev, 0.000005)
toc = time.time()
print('Naive loss and gradient: computed in %fs' % (toc - tic))

tic = time.time()
_, grad_vectorized = svm_loss_vectorized(W, X_dev, y_dev, 0.000005)
toc = time.time()
print('Vectorized loss and gradient: computed in %fs' % (toc - tic))

# The loss is a single number, so it is easy to compare the values computed
# by the two implementations. The gradient on the other hand is a matrix, so
# we use the Frobenius norm to compare them.
difference = np.linalg.norm(grad_naive - grad_vectorized, ord='fro')
print('difference: %f' % difference)
```

Naive loss and gradient: computed in 0.126098s
 Vectorized loss and gradient: computed in 0.010333s
 difference: 0.000000

1.2.1 Stochastic Gradient Descent

We now have vectorized and efficient expressions for the loss, the gradient and our gradient matches the numerical gradient. We are therefore ready to do SGD to minimize the loss. Your code for this part will be written inside `cs231n/classifiers/linear_classifier.py`.


[30]: *# In the file linear_classifier.py, implement SGD in the function*
LinearClassifier.train() and then run it with the code below.

```
from cs231n.classifiers import LinearSVM
svm = LinearSVM()
tic = time.time()
loss_hist = svm.train(X_train, y_train, learning_rate=1e-7, reg=2.5e4,
                      num_iters=1500, verbose=True)
toc = time.time()
```

```
print('That took %fs' % (toc - tic))
```

```
iteration 0 / 1500: loss 793.317377
iteration 100 / 1500: loss 288.854060
iteration 200 / 1500: loss 108.680140
iteration 300 / 1500: loss 42.577876
iteration 400 / 1500: loss 19.192679
iteration 500 / 1500: loss 10.578673
iteration 600 / 1500: loss 7.403988
iteration 700 / 1500: loss 5.916690
iteration 800 / 1500: loss 5.139422
iteration 900 / 1500: loss 5.381073
iteration 1000 / 1500: loss 4.819770
iteration 1100 / 1500: loss 5.234745
iteration 1200 / 1500: loss 4.931539
iteration 1300 / 1500: loss 5.221948
iteration 1400 / 1500: loss 5.024407
That took 7.548413s
```

```
[31]: # A useful debugging strategy is to plot the loss as a function of
      # iteration number:
      plt.plot(loss_hist)
      plt.xlabel('Iteration number')
      plt.ylabel('Loss value')
      plt.show()
```



```
[32]: # Write the LinearSVM.predict function and evaluate the performance on both the
# training and validation set
y_train_pred = svm.predict(X_train)
print('training accuracy: %f' % (np.mean(y_train == y_train_pred), ))
y_val_pred = svm.predict(X_val)
print('validation accuracy: %f' % (np.mean(y_val == y_val_pred), ))
```

```
training accuracy: 0.371531
validation accuracy: 0.386000
```

```
[33]: # Use the validation set to tune hyperparameters (regularization strength and
# learning rate). You should experiment with different ranges for the learning
# rates and regularization strengths; if you are careful you should be able to
# get a classification accuracy of about 0.39 (> 0.385) on the validation set.

# Note: you may see runtime/overflow warnings during hyper-parameter search.
# This may be caused by extreme values, and is not a bug.

# results is dictionary mapping tuples of the form
# (learning_rate, regularization_strength) to tuples of the form
```

```

# (training_accuracy, validation_accuracy). The accuracy is simply the fraction
# of data points that are correctly classified.
results = {}
best_val = -1 # The highest validation accuracy that we have seen so far.
best_svm = None # The LinearSVM object that achieved the highest validation
    ↳rate.

#####
# TODO:
# Write code that chooses the best hyperparameters by tuning on the validation #
# set. For each combination of hyperparameters, train a linear SVM on the #
# training set, compute its accuracy on the training and validation sets, and #
# store these numbers in the results dictionary. In addition, store the best #
# validation accuracy in best_val and the LinearSVM object that achieves this #
# accuracy in best_svm.
#
# Hint: You should use a small value for num_iters as you develop your #
# validation code so that the SVMs don't take much time to train; once you are #
# confident that your validation code works, you should rerun the validation #
# code with a larger value for num_iters.
#####

# Provided as a reference. You may or may not want to change these
    ↳hyperparameters
learning_rates = [0.85e-7, 0.9e-7, 0.95e-7, 1e-7]
regularization_strengths = [5.5e4, 6e4]

#best validation accuracy obtained for
#lr 9.000000e-08 reg 6.000000e+04 train accuracy: 0.369551 val accuracy: 0.
    ↳391000

# *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
for lr in learning_rates:
    for reg in regularization_strengths:
        curr_svm = LinearSVM()
        loss_hist = curr_svm.train(X_train, y_train, learning_rate=lr, reg=reg,
                                   num_iters=1500, verbose=False)

        y_train_pred = svm.predict(X_train)
        train_accuracy = np.mean(y_train == y_train_pred)

        y_val_pred = svm.predict(X_val)
        val_accuracy = np.mean(y_val == y_val_pred)
        if (val_accuracy > best_val):
            best_val = val_accuracy
            best_svm = curr_svm

```



```

        results[(lr, reg)] = (train_accuracy, val_accuracy)

# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

# Print out results.
for lr, reg in sorted(results):
    train_accuracy, val_accuracy = results[(lr, reg)]
    print('lr %e reg %e train accuracy: %f val accuracy: %f' % (
        lr, reg, train_accuracy, val_accuracy))

print('best validation accuracy achieved during cross-validation: %f' %
      ↪best_val)

```

```

lr 8.500000e-08 reg 5.500000e+04 train accuracy: 0.371531 val accuracy: 0.386000
lr 8.500000e-08 reg 6.000000e+04 train accuracy: 0.371531 val accuracy: 0.386000
lr 9.000000e-08 reg 5.500000e+04 train accuracy: 0.371531 val accuracy: 0.386000
lr 9.000000e-08 reg 6.000000e+04 train accuracy: 0.371531 val accuracy: 0.386000
lr 9.500000e-08 reg 5.500000e+04 train accuracy: 0.371531 val accuracy: 0.386000
lr 9.500000e-08 reg 6.000000e+04 train accuracy: 0.371531 val accuracy: 0.386000
lr 1.000000e-07 reg 5.500000e+04 train accuracy: 0.371531 val accuracy: 0.386000
lr 1.000000e-07 reg 6.000000e+04 train accuracy: 0.371531 val accuracy: 0.386000
best validation accuracy achieved during cross-validation: 0.386000

```

```

[34]: # Visualize the cross-validation results
import math
import pdb

# pdb.set_trace()

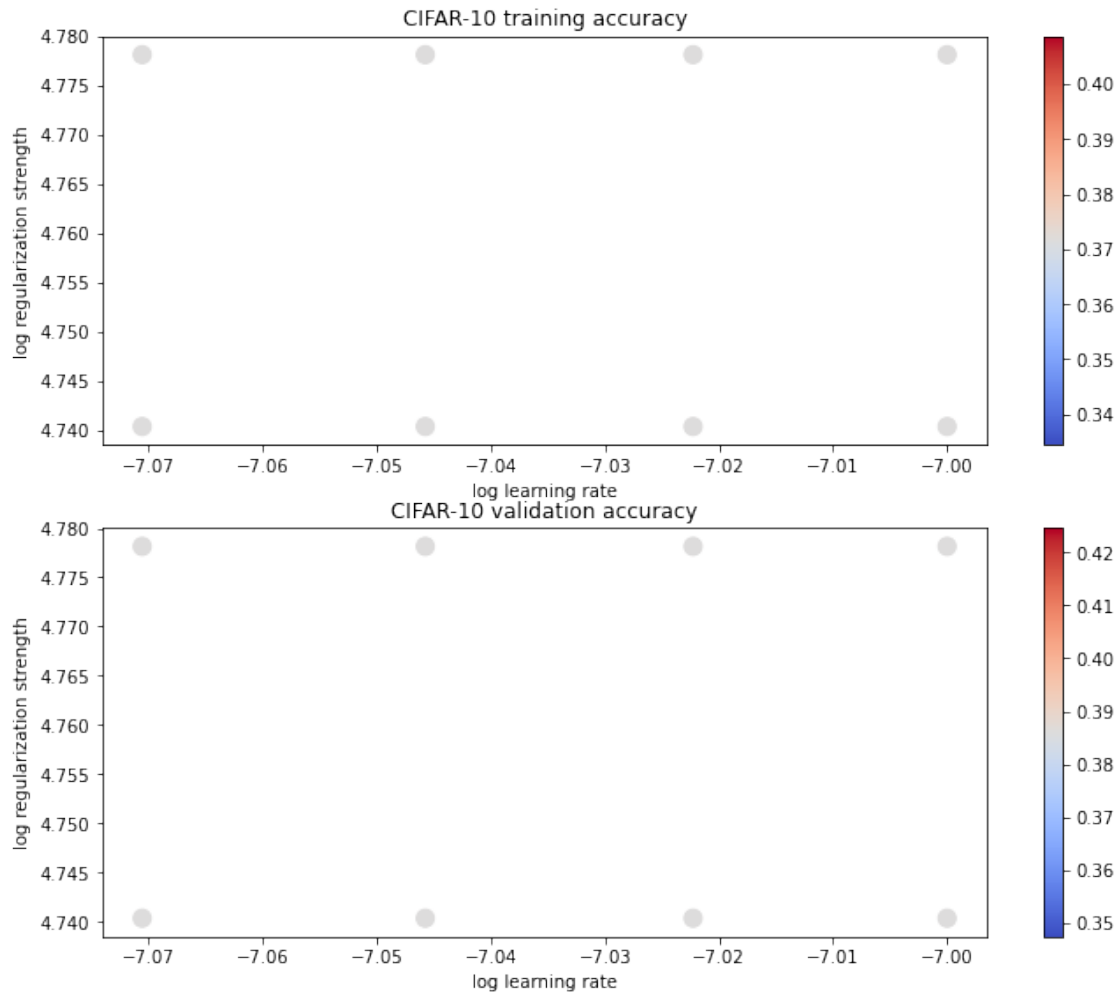
x_scatter = [math.log10(x[0]) for x in results]
y_scatter = [math.log10(x[1]) for x in results]

# plot training accuracy
marker_size = 100
colors = [results[x][0] for x in results]
plt.subplot(2, 1, 1)
plt.tight_layout(pad=3)
plt.scatter(x_scatter, y_scatter, marker_size, c=colors, cmap=plt.cm.coolwarm)
plt.colorbar()
plt.xlabel('log learning rate')
plt.ylabel('log regularization strength')
plt.title('CIFAR-10 training accuracy')

# plot validation accuracy
colors = [results[x][1] for x in results] # default size of markers is 20
plt.subplot(2, 1, 2)

```

```
plt.scatter(x_scatter, y_scatter, marker_size, c=colors, cmap=plt.cm.coolwarm)
plt.colorbar()
plt.xlabel('log learning rate')
plt.ylabel('log regularization strength')
plt.title('CIFAR-10 validation accuracy')
plt.show()
```



```
[35]: # Evaluate the best svm on test set
y_test_pred = best_svm.predict(X_test)
test_accuracy = np.mean(y_test == y_test_pred)
print('linear SVM on raw pixels final test set accuracy: %f' % test_accuracy)
```

linear SVM on raw pixels final test set accuracy: 0.365000

```
[36]: # Visualize the learned weights for each class.
```

```

# Depending on your choice of learning rate and regularization strength, these
→ may
# or may not be nice to look at.
w = best_svm.W[:-1,:] # strip out the bias
w = w.reshape(32, 32, 3, 10)
w_min, w_max = np.min(w), np.max(w)
classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', '
→ 'ship', 'truck']
for i in range(10):
    plt.subplot(2, 5, i + 1)

    # Rescale the weights to be between 0 and 255
    wimg = 255.0 * (w[:, :, :, i].squeeze() - w_min) / (w_max - w_min)
    plt.imshow(wimg.astype('uint8'))
    plt.axis('off')
    plt.title(classes[i])

```



Inline question 2

Describe what your visualized SVM weights look like, and offer a brief explanation for why they look the way they do.

Your Answer : Visualized SVM weights represents templates for each of the class (looks similar to lecture slide). \ Interpretation: \ Ship, airplane has majority of the blue pixels indicaitng majority of the ship and airplane having background of water and sky respectively.

Weights of horse indicated horse standing in left and right direction, as an average of several images which contains horse facing left or right.

Frog has majority of the pixels with green color in center indicating the color of frog being green.

\

softmax

April 16, 2022

```
[1]: # This mounts your Google Drive to the Colab VM.
from google.colab import drive
drive.mount('/content/drive')

# TODO: Enter the foldername in your Drive where you have saved the unzipped
# assignment folder, e.g. 'cs231n/assignments/assignment1/'
FOLDERNAME = 'Colab_Notebooks/cs231n/assignments/assignment1/'
assert FOLDERNAME is not None, "[!] Enter the foldername."

# Now that we've mounted your Drive, this ensures that
# the Python interpreter of the Colab VM can load
# python files from within it.
import sys
sys.path.append('/content/drive/My Drive/{}'.format(FOLDERNAME))

# This downloads the CIFAR-10 dataset to your Drive
# if it doesn't already exist.
%cd /content/drive/My\ Drive/$FOLDERNAME/cs231n/datasets/
!bash get_datasets.sh
%cd /content/drive/My\ Drive/$FOLDERNAME
```

```
Mounted at /content/drive
/content/drive/My
Drive/Colab_Notebooks/cs231n/assignments/assignment1/cs231n/datasets
/content/drive/My Drive/Colab_Notebooks/cs231n/assignments/assignment1
```

1 Softmax exercise

Complete and hand in this completed worksheet (including its outputs and any supporting code outside of the worksheet) with your assignment submission. For more details see the [assignments page](#) on the course website.

This exercise is analogous to the SVM exercise. You will:

- implement a fully-vectorized **loss function** for the Softmax classifier
- implement the fully-vectorized expression for its **analytic gradient**
- **check your implementation** with numerical gradient

- use a validation set to **tune the learning rate and regularization** strength
- **optimize** the loss function with **SGD**
- **visualize** the final learned weights

```
[2]: import random
import numpy as np
from cs231n.data_utils import load_CIFAR10
import matplotlib.pyplot as plt

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/
→ autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

[3]: def get_CIFAR10_data(num_training=49000, num_validation=1000, num_test=1000,
→ num_dev=500):
    """
    Load the CIFAR-10 dataset from disk and perform preprocessing to prepare
    it for the linear classifier. These are the same steps as we used for the
    SVM, but condensed to a single function.
    """
    # Load the raw CIFAR-10 data
    cifar10_dir = 'cs231n/datasets/cifar-10-batches-py'

    # Cleaning up variables to prevent loading data multiple times (which may
    → cause memory issue)
    try:
        del X_train, y_train
        del X_test, y_test
        print('Clear previously loaded data.')
    except:
        pass

    X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)

    # subsample the data
    mask = list(range(num_training, num_training + num_validation))
    X_val = X_train[mask]
    y_val = y_train[mask]
    mask = list(range(num_training))
    X_train = X_train[mask]
```

```

y_train = y_train[mask]
mask = list(range(num_test))
X_test = X_test[mask]
y_test = y_test[mask]
mask = np.random.choice(num_training, num_dev, replace=False)
X_dev = X_train[mask]
y_dev = y_train[mask]

# Preprocessing: reshape the image data into rows
X_train = np.reshape(X_train, (X_train.shape[0], -1))
X_val = np.reshape(X_val, (X_val.shape[0], -1))
X_test = np.reshape(X_test, (X_test.shape[0], -1))
X_dev = np.reshape(X_dev, (X_dev.shape[0], -1))

# Normalize the data: subtract the mean image
mean_image = np.mean(X_train, axis = 0)
X_train -= mean_image
X_val -= mean_image
X_test -= mean_image
X_dev -= mean_image

# add bias dimension and transform into columns
X_train = np.hstack([X_train, np.ones((X_train.shape[0], 1))])
X_val = np.hstack([X_val, np.ones((X_val.shape[0], 1))])
X_test = np.hstack([X_test, np.ones((X_test.shape[0], 1))])
X_dev = np.hstack([X_dev, np.ones((X_dev.shape[0], 1))])

return X_train, y_train, X_val, y_val, X_test, y_test, X_dev, y_dev

# Invoke the above function to get our data.
X_train, y_train, X_val, y_val, X_test, y_test, X_dev, y_dev = ↳
get_CIFAR10_data()
print('Train data shape: ', X_train.shape)
print('Train labels shape: ', y_train.shape)
print('Validation data shape: ', X_val.shape)
print('Validation labels shape: ', y_val.shape)
print('Test data shape: ', X_test.shape)
print('Test labels shape: ', y_test.shape)
print('dev data shape: ', X_dev.shape)
print('dev labels shape: ', y_dev.shape)

```

```

Train data shape: (49000, 3073)
Train labels shape: (49000,)
Validation data shape: (1000, 3073)
Validation labels shape: (1000,)
Test data shape: (1000, 3073)

```

```
Test labels shape: (1000,)
dev data shape: (500, 3073)
dev labels shape: (500,)
```

1.1 Softmax Classifier

Your code for this section will all be written inside `cs231n/classifiers/softmax.py`.

```
[4]: # First implement the naive softmax loss function with nested loops.
# Open the file cs231n/classifiers/softmax.py and implement the
# softmax_loss_naive function.

from cs231n.classifiers.softmax import softmax_loss_naive
import time

# Generate a random softmax weight matrix and use it to compute the loss.
W = np.random.randn(3073, 10) * 0.0001
loss, grad = softmax_loss_naive(W, X_dev, y_dev, 0.0)

# As a rough sanity check, our loss should be something close to -log(0.1).
print('loss: %f' % loss)
print('sanity check: %f' % (-np.log(0.1)))
```

```
loss: 2.334037
sanity check: 2.302585
```

Inline Question 1

Why do we expect our loss to be close to $-\log(0.1)$? Explain briefly.

Your Answer : \ As mention in class it is because \ $-\log(P(y) * (P(y) / Q(y))) = -\log(1/c / c*c) = -\log(1/c)$. For our case since $c = 10$, \ hence we have $-\log(1/10) = -\log(0.1) = 2.302585$

```
[5]: # Complete the implementation of softmax_loss_naive and implement a (naive)
# version of the gradient that uses nested loops.
loss, grad = softmax_loss_naive(W, X_dev, y_dev, 0.0)

# As we did for the SVM, use numeric gradient checking as a debugging tool.
# The numeric gradient should be close to the analytic gradient.
from cs231n.gradient_check import grad_check_sparse
f = lambda w: softmax_loss_naive(w, X_dev, y_dev, 0.0)[0]
grad_numerical = grad_check_sparse(f, W, grad, 10)

# similar to SVM case, do another gradient check with regularization
loss, grad = softmax_loss_naive(W, X_dev, y_dev, 5e1)
f = lambda w: softmax_loss_naive(w, X_dev, y_dev, 5e1)[0]
grad_numerical = grad_check_sparse(f, W, grad, 10)
```



```

numerical: -0.287850 analytic: -0.287850, relative error: 1.617901e-07
numerical: -4.095274 analytic: -4.095274, relative error: 1.124373e-08
numerical: -0.275564 analytic: -0.275564, relative error: 6.636527e-08
numerical: -0.863496 analytic: -0.863496, relative error: 1.555851e-08
numerical: -1.519414 analytic: -1.519414, relative error: 1.311175e-08
numerical: 0.875556 analytic: 0.875556, relative error: 5.701934e-09
numerical: -3.152644 analytic: -3.152644, relative error: 3.473895e-09
numerical: 3.421196 analytic: 3.421196, relative error: 6.682175e-09
numerical: 1.181829 analytic: 1.181828, relative error: 2.661630e-09
numerical: 0.520562 analytic: 0.520562, relative error: 8.666974e-08
numerical: 0.845485 analytic: 0.845485, relative error: 2.274206e-08
numerical: -1.520918 analytic: -1.520918, relative error: 8.004306e-09
numerical: -1.827895 analytic: -1.827895, relative error: 2.483346e-09
numerical: -1.844154 analytic: -1.844154, relative error: 3.407018e-08
numerical: -0.052893 analytic: -0.052893, relative error: 5.565774e-07
numerical: 3.224674 analytic: 3.224674, relative error: 6.708838e-10
numerical: 2.499276 analytic: 2.499276, relative error: 1.745426e-08
numerical: -1.172722 analytic: -1.172722, relative error: 2.497786e-08
numerical: -1.059227 analytic: -1.059227, relative error: 1.642319e-09
numerical: 4.432659 analytic: 4.432659, relative error: 4.406464e-09

```

```

[6]: # Now that we have a naive implementation of the softmax loss function and its
      →gradient,
      # implement a vectorized version in softmax_loss_vectorized.
      # The two versions should compute the same results, but the vectorized version
      →should be
      # much faster.
      #https://mlxai.github.io/2017/01/09/
      →implementing-softmax-classifier-with-vectorized-operations.html
      tic = time.time()
      loss_naive, grad_naive = softmax_loss_naive(W, X_dev, y_dev, 0.000005)
      toc = time.time()
      print('naive loss: %e computed in %fs' % (loss_naive, toc - tic))

      from cs231n.classifiers.softmax import softmax_loss_vectorized
      tic = time.time()
      loss_vectorized, grad_vectorized = softmax_loss_vectorized(W, X_dev, y_dev, 0.
      →000005)
      toc = time.time()
      print('vectorized loss: %e computed in %fs' % (loss_vectorized, toc - tic))

      # As we did for the SVM, we use the Frobenius norm to compare the two versions
      # of the gradient.
      grad_difference = np.linalg.norm(grad_naive - grad_vectorized, ord='fro')
      print('Loss difference: %f' % np.abs(loss_naive - loss_vectorized))
      print('Gradient difference: %f' % grad_difference)

```

naive loss: 2.334037e+00 computed in 0.119939s
vectorized loss: 2.334037e+00 computed in 0.017927s
Loss difference: 0.000000
Gradient difference: 0.000000

```
[7]: # Use the validation set to tune hyperparameters (regularization strength and
# learning rate). You should experiment with different ranges for the learning
# rates and regularization strengths; if you are careful you should be able to
# get a classification accuracy of over 0.35 on the validation set.

from cs231n.classifiers import Softmax
results = {}
best_val = -1
best_softmax = None

#####
# TODO:
# Use the validation set to set the learning rate and regularization strength. #
# This should be identical to the validation that you did for the SVM; save   #
# the best trained softmax classifier in best_softmax.                         #
#####

# Provided as a reference. You may or may not want to change these
# hyperparameters
learning_rates = [2.4e-7, 2.5e-7, 2.8e-7, 3e-7] #1e-7, 6e-7
regularization_strengths = [2.4e4, 2.5e4, 2.8e4] #5e4

# *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
for lr in learning_rates:
    for reg in regularization_strengths:
        curr_softmax = Softmax()
        loss_hist = curr_softmax.train(X_train, y_train, learning_rate=lr, reg=reg,
                                       num_iters=1000, verbose=False)

        y_train_pred = curr_softmax.predict(X_train)
        train_accuracy = np.mean(y_train == y_train_pred)

        y_val_pred = curr_softmax.predict(X_val)
        val_accuracy = np.mean(y_val == y_val_pred)
        if (val_accuracy > best_val):
            best_val = val_accuracy
            best_softmax = curr_softmax

    results[(lr, reg)] = (train_accuracy, val_accuracy)
```

```
# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

# Print out results.
for lr, reg in sorted(results):
    train_accuracy, val_accuracy = results[(lr, reg)]
    print('lr %e reg %e train accuracy: %f val accuracy: %f' % (
        lr, reg, train_accuracy, val_accuracy))

print('best validation accuracy achieved during cross-validation: %f' %
      ↪best_val)
```

```
lr 2.400000e-07 reg 2.400000e+04 train accuracy: 0.348571 val accuracy: 0.350000
lr 2.400000e-07 reg 2.500000e+04 train accuracy: 0.344898 val accuracy: 0.359000
lr 2.400000e-07 reg 2.800000e+04 train accuracy: 0.350857 val accuracy: 0.358000
lr 2.500000e-07 reg 2.400000e+04 train accuracy: 0.350490 val accuracy: 0.369000
lr 2.500000e-07 reg 2.500000e+04 train accuracy: 0.355490 val accuracy: 0.360000
lr 2.500000e-07 reg 2.800000e+04 train accuracy: 0.346551 val accuracy: 0.369000
lr 2.800000e-07 reg 2.400000e+04 train accuracy: 0.351694 val accuracy: 0.357000
lr 2.800000e-07 reg 2.500000e+04 train accuracy: 0.352102 val accuracy: 0.369000
lr 2.800000e-07 reg 2.800000e+04 train accuracy: 0.336551 val accuracy: 0.352000
lr 3.000000e-07 reg 2.400000e+04 train accuracy: 0.351592 val accuracy: 0.361000
lr 3.000000e-07 reg 2.500000e+04 train accuracy: 0.348592 val accuracy: 0.356000
lr 3.000000e-07 reg 2.800000e+04 train accuracy: 0.344980 val accuracy: 0.348000
best validation accuracy achieved during cross-validation: 0.369000
```

```
[8]: # evaluate on test set
# Evaluate the best softmax on test set
y_test_pred = best_softmax.predict(X_test)
test_accuracy = np.mean(y_test == y_test_pred)
print('softmax on raw pixels final test set accuracy: %f' % (test_accuracy, ))
```

softmax on raw pixels final test set accuracy: 0.360000

Inline Question 2 - True or False

Suppose the overall training loss is defined as the sum of the per-datapoint loss over all training examples. It is possible to add a new datapoint to a training set that would leave the SVM loss unchanged, but this is not the case with the Softmax classifier loss.

Your Answer : True

Your Explanation : Training loss for a given training example in SVM is considered if $\text{margin} > 0$. If the $\text{margin} = 0$, SVM loss is not considered for that particular training example, hence in the event when an added data point has $\text{margin} = 0$, training loss remains unchanged. \ However, softmax loss of a given training example = $-\text{np.log}(P(\text{correct_class}))$. Predicted probability for the correct class can never be equal to 1 because softmax prediction for multiple classes are distributed between all classes and adds upto 1. Hence it is not possible to have 0 loss for softmax.

```
[9]: # Visualize the learned weights for each class
w = best_softmax.W[:-1,:] # strip out the bias
w = w.reshape(32, 32, 3, 10)

w_min, w_max = np.min(w), np.max(w)

classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
for i in range(10):
    plt.subplot(2, 5, i + 1)

    # Rescale the weights to be between 0 and 255
    wimg = 255.0 * (w[:, :, :, i].squeeze() - w_min) / (w_max - w_min)
    plt.imshow(wimg.astype('uint8'))
    plt.axis('off')
    plt.title(classes[i])
```



two_layer_net

April 16, 2022

```
[1]: # This mounts your Google Drive to the Colab VM.
from google.colab import drive
drive.mount('/content/drive')

# TODO: Enter the foldername in your Drive where you have saved the unzipped
# assignment folder, e.g. 'cs231n/assignments/assignment1/'
FOLDERNAME = 'Colab_Notebooks/cs231n/assignments/assignment1/'
assert FOLDERNAME is not None, "[!] Enter the foldername."

# Now that we've mounted your Drive, this ensures that
# the Python interpreter of the Colab VM can load
# python files from within it.
import sys
sys.path.append('/content/drive/My Drive/{}'.format(FOLDERNAME))

# This downloads the CIFAR-10 dataset to your Drive
# if it doesn't already exist.
%cd /content/drive/My\ Drive/$FOLDERNAME/cs231n/datasets/
!bash get_datasets.sh
%cd /content/drive/My\ Drive/$FOLDERNAME
```

```
Mounted at /content/drive
/content/drive/My
Drive/Colab_Notebooks/cs231n/assignments/assignment1/cs231n/datasets
/content/drive/My Drive/Colab_Notebooks/cs231n/assignments/assignment1
```

1 Fully-Connected Neural Nets

In this exercise we will implement fully-connected networks using a modular approach. For each layer we will implement a `forward` and a `backward` function. The `forward` function will receive inputs, weights, and other parameters and will return both an output and a `cache` object storing data needed for the backward pass, like this:

```
def layer_forward(x, w):
    """ Receive inputs x and weights w """
    # Do some computations ...
    z = # ... some intermediate value
```

```

# Do some more computations ...
out = # the output

cache = (x, w, z, out) # Values we need to compute gradients

return out, cache

```

The backward pass will receive upstream derivatives and the `cache` object, and will return gradients with respect to the inputs and weights, like this:

```

def layer_backward(dout, cache):
    """
    Receive dout (derivative of loss with respect to outputs) and cache,
    and compute derivative with respect to inputs.
    """
    # Unpack cache values
    x, w, z, out = cache

    # Use values in cache to compute derivatives
    dx = # Derivative of loss with respect to x
    dw = # Derivative of loss with respect to w

    return dx, dw

```

After implementing a bunch of layers this way, we will be able to easily combine them to build classifiers with different architectures.

```

[2]: # As usual, a bit of setup
from __future__ import print_function
import time
import numpy as np
import matplotlib.pyplot as plt
from cs231n.classifiers.fc_net import *
from cs231n.data_utils import get_CIFAR10_data
from cs231n.gradient_check import eval_numerical_gradient, eval_numerical_gradient_array
from cs231n.solver import Solver

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/
→autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

```

```
def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))
```

[3]: *# Load the (preprocessed) CIFAR10 data.*

```
data = get_CIFAR10_data()
for k, v in list(data.items()):
    print('%s: ' % k, v.shape)
```

```
('X_train: ', (49000, 3, 32, 32))
('y_train: ', (49000,))
('X_val: ', (1000, 3, 32, 32))
('y_val: ', (1000,))
('X_test: ', (1000, 3, 32, 32))
('y_test: ', (1000,))
```

2 Affine layer: forward

Open the file `cs231n/layers.py` and implement the `affine_forward` function.

Once you are done you can test your implementation by running the following:

[4]: *# Test the affine_forward function*

```
num_inputs = 2
input_shape = (4, 5, 6)
output_dim = 3

input_size = num_inputs * np.prod(input_shape)
weight_size = output_dim * np.prod(input_shape)

x = np.linspace(-0.1, 0.5, num=input_size).reshape(num_inputs, *input_shape)
w = np.linspace(-0.2, 0.3, num=weight_size).reshape(np.prod(input_shape),
→output_dim)
b = np.linspace(-0.3, 0.1, num=output_dim)

out, _ = affine_forward(x, w, b)
correct_out = np.array([[ 1.49834967,  1.70660132,  1.91485297],
                        [ 3.25553199,  3.5141327,  3.77273342]])

# Compare your output with ours. The error should be around e-9 or less.
print('Testing affine_forward function:')
print('difference: ', rel_error(out, correct_out))
```

```
Testing affine_forward function:
difference: 9.769849468192957e-10
```

3 Affine layer: backward

Now implement the `affine_backward` function and test your implementation using numeric gradient checking.

```
[5]: # Test the affine_backward function
np.random.seed(231)
x = np.random.randn(10, 2, 3)
w = np.random.randn(6, 5)
b = np.random.randn(5)
dout = np.random.randn(10, 5)

dx_num = eval_numerical_gradient_array(lambda x: affine_forward(x, w, b)[0], x,
    ↪dout)
dw_num = eval_numerical_gradient_array(lambda w: affine_forward(x, w, b)[0], w,
    ↪dout)
db_num = eval_numerical_gradient_array(lambda b: affine_forward(x, w, b)[0], b,
    ↪dout)

_, cache = affine_forward(x, w, b)
dx, dw, db = affine_backward(dout, cache)

# The error should be around e-10 or less
print('Testing affine_backward function:')
print('dx error: ', rel_error(dx_num, dx))
print('dw error: ', rel_error(dw_num, dw))
print('db error: ', rel_error(db_num, db))
```

```
Testing affine_backward function:
dx error:  5.399100368651805e-11
dw error:  9.904211865398145e-11
db error:  2.4122867568119087e-11
```

4 ReLU activation: forward

Implement the forward pass for the ReLU activation function in the `relu_forward` function and test your implementation using the following:

```
[6]: # Test the relu_forward function

x = np.linspace(-0.5, 0.5, num=12).reshape(3, 4)

out, _ = relu_forward(x)
correct_out = np.array([[ 0.,          0.,          0.,          0.],
                        [ 0.,          0.,          0.04545455, 0.13636364],
                        [ 0.22727273, 0.31818182, 0.40909091, 0.5]])
```



```
# Compare your output with ours. The error should be on the order of e-8
print('Testing relu_forward function:')
print('difference: ', rel_error(out, correct_out))
```

```
Testing relu_forward function:
difference: 4.999999798022158e-08
```

5 ReLU activation: backward

Now implement the backward pass for the ReLU activation function in the `relu_backward` function and test your implementation using numeric gradient checking:

```
[7]: np.random.seed(231)
x = np.random.randn(10, 10)
dout = np.random.randn(*x.shape)

dx_num = eval_numerical_gradient_array(lambda x: relu_forward(x)[0], x, dout)

_, cache = relu_forward(x)
dx = relu_backward(dout, cache)

# The error should be on the order of e-12
print('Testing relu_backward function:')
print('dx error: ', rel_error(dx_num, dx))
```

```
Testing relu_backward function:
dx error: 3.2756349136310288e-12
```

5.1 Inline Question 1:

We've only asked you to implement ReLU, but there are a number of different activation functions that one could use in neural networks, each with its pros and cons. In particular, an issue commonly seen with activation functions is getting zero (or close to zero) gradient flow during backpropagation. Which of the following activation functions have this problem? If you consider these functions in the one dimensional case, what types of input would lead to this behaviour? 1. Sigmoid 2. ReLU 3. Leaky ReLU

5.2 Answer:

1. Sigmoid faces the issue of vanishing gradient (derivative close to zero), once the function approaches 1 (for $x > 0$) or 0 (for $x < 0$). For inputs in 1D, for as $\text{sigmoid}(x)$ is ~ 1 or ~ 0 , the derivative tends to become 0 with minimal slope.
2. ReLU has 0 derivative for $x < 0$ and derivative = 1 for all $x > 0$. For 1D inputs, ReLU would have the issue of zero gradient when $x < 0$.

Leaky ReLU does NOT encounter gradient approaching zero, because $f'(x) = 1$ (for $x > 0$) and $f'(x) = a$ (constant value) (for $x < 0$)

6 "Sandwich" layers

There are some common patterns of layers that are frequently used in neural nets. For example, affine layers are frequently followed by a ReLU nonlinearity. To make these common patterns easy, we define several convenience layers in the file `cs231n/layer_utils.py`.

For now take a look at the `affine_relu_forward` and `affine_relu_backward` functions, and run the following to numerically gradient check the backward pass:

```
[8]: from cs231n.layer_utils import affine_relu_forward, affine_relu_backward
np.random.seed(231)
x = np.random.randn(2, 3, 4)
w = np.random.randn(12, 10)
b = np.random.randn(10)
dout = np.random.randn(2, 10)

out, cache = affine_relu_forward(x, w, b)
dx, dw, db = affine_relu_backward(dout, cache)

dx_num = eval_numerical_gradient_array(lambda x: affine_relu_forward(x, w,
    ↪b)[0], x, dout)
dw_num = eval_numerical_gradient_array(lambda w: affine_relu_forward(x, w,
    ↪b)[0], w, dout)
db_num = eval_numerical_gradient_array(lambda b: affine_relu_forward(x, w,
    ↪b)[0], b, dout)

# Relative error should be around e-10 or less
print('Testing affine_relu_forward and affine_relu_backward:')
print('dx error: ', rel_error(dx_num, dx))
print('dw error: ', rel_error(dw_num, dw))
print('db error: ', rel_error(db_num, db))
```

Testing affine_relu_forward and affine_relu_backward:

dx error: 2.299579177309368e-11

dw error: 8.162011105764925e-11

db error: 7.826724021458994e-12

7 Loss layers: Softmax and SVM

Now implement the loss and gradient for softmax and SVM in the `softmax_loss` and `svm_loss` function in `cs231n/layers.py`. These should be similar to what you implemented in `cs231n/classifiers/softmax.py` and `cs231n/classifiers/linear_svm.py`.

You can make sure that the implementations are correct by running the following:

```
[9]: np.random.seed(231)
num_classes, num_inputs = 10, 50
x = 0.001 * np.random.randn(num_inputs, num_classes)
y = np.random.randint(num_classes, size=num_inputs)

dx_num = eval_numerical_gradient(lambda x: svm_loss(x, y)[0], x, verbose=False)
loss, dx = svm_loss(x, y)

# Test svm_loss function. Loss should be around 9 and dx error should be around
# the order of e-9
print('Testing svm_loss:')
print('loss: ', loss)
print('dx error: ', rel_error(dx_num, dx))

dx_num = eval_numerical_gradient(lambda x: softmax_loss(x, y)[0], x,
# verbose=False)
loss, dx = softmax_loss(x, y)

# Test softmax_loss function. Loss should be close to 2.3 and dx error should
# be around e-8
print('\nTesting softmax_loss:')
print('loss: ', loss)
print('dx error: ', rel_error(dx_num, dx))
```

```
Testing svm_loss:
loss: 8.999602749096233
dx error: 1.4021566006651672e-09
```

```
Testing softmax_loss:
loss: 2.3025458445007376
dx error: 8.234144091578429e-09
```

8 Two-layer network

Open the file `cs231n/classifiers/fc_net.py` and complete the implementation of the `TwoLayerNet` class. Read through it to make sure you understand the API. You can run the cell below to test your implementation.

```
[10]: np.random.seed(231)
N, D, H, C = 3, 5, 50, 7
X = np.random.randn(N, D)
y = np.random.randint(C, size=N)

std = 1e-3
```

```

model = TwoLayerNet(input_dim=D, hidden_dim=H, num_classes=C, weight_scale=std)

print('Testing initialization ... ')
W1_std = abs(model.params['W1'].std() - std)
b1 = model.params['b1']
W2_std = abs(model.params['W2'].std() - std)
b2 = model.params['b2']
assert W1_std < std / 10, 'First layer weights do not seem right'
assert np.all(b1 == 0), 'First layer biases do not seem right'
assert W2_std < std / 10, 'Second layer weights do not seem right'
assert np.all(b2 == 0), 'Second layer biases do not seem right'

print('Testing test-time forward pass ... ')
model.params['W1'] = np.linspace(-0.7, 0.3, num=D*H).reshape(D, H)
model.params['b1'] = np.linspace(-0.1, 0.9, num=H)
model.params['W2'] = np.linspace(-0.3, 0.4, num=H*C).reshape(H, C)
model.params['b2'] = np.linspace(-0.9, 0.1, num=C)
X = np.linspace(-5.5, 4.5, num=N*D).reshape(D, N).T
scores = model.loss(X)
correct_scores = np.asarray(
    [[11.53165108, 12.2917344, 13.05181771, 13.81190102, 14.57198434, 15.
    ↪33206765, 16.09215096],
     [12.05769098, 12.74614105, 13.43459113, 14.1230412, 14.81149128, 15.
    ↪49994135, 16.18839143],
     [12.58373087, 13.20054771, 13.81736455, 14.43418138, 15.05099822, 15.
    ↪66781506, 16.2846319 ]])
scores_diff = np.abs(scores - correct_scores).sum()

assert scores_diff < 1e-6, 'Problem with test-time forward pass'

print('Testing training loss (no regularization)')
y = np.asarray([0, 5, 1])
loss, grads = model.loss(X, y)
#print(loss)
correct_loss = 3.4702243556
assert abs(loss - correct_loss) < 1e-10, 'Problem with training-time loss'

model.reg = 1.0
loss, grads = model.loss(X, y)
print(loss)
correct_loss = 26.5948426952
assert abs(loss - correct_loss) < 1e-10, 'Problem with regularization loss'

# Errors should be around e-7 or less
for reg in [0.0, 0.7]:
    print('Running numeric gradient check with reg = ', reg)
    model.reg = reg

```

```

loss, grads = model.loss(X, y)

for name in sorted(grads):
    f = lambda _: model.loss(X, y)[0]
    grad_num = eval_numerical_gradient(f, model.params[name], verbose=False)
    print('%s relative error: %.2e' % (name, rel_error(grad_num, grads[name])))

```

```

Testing initialization ...
Testing test-time forward pass ...
Testing training loss (no regularization)
26.594842695238583
Running numeric gradient check with reg = 0.0
W1 relative error: 1.83e-08
W2 relative error: 3.20e-10
b1 relative error: 9.83e-09
b2 relative error: 4.33e-10
Running numeric gradient check with reg = 0.7
W1 relative error: 2.53e-07
W2 relative error: 2.85e-08
b1 relative error: 1.56e-08
b2 relative error: 9.09e-10

```

9 Solver

Open the file `cs231n/solver.py` and read through it to familiarize yourself with the API. You also need to implement the `sgd` function in `cs231n/optim.py`. After doing so, use a `Solver` instance to train a `TwoLayerNet` that achieves about 36% accuracy on the validation set.

```

[11]: input_size = 32 * 32 * 3
hidden_size = 50
num_classes = 10
model = TwoLayerNet(input_size, hidden_size, num_classes)
solver = None

#####
# TODO: Use a Solver instance to train a TwoLayerNet that achieves about 36% #
# accuracy on the validation set.                                           #
#####
# *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
solver = Solver(model, data,
                update_rule = 'sgd',
                optim_config = {
                    "learning_rate": 1e-3,
                },
                lr_decay = 0.95,
                num_epochs = 5, batch_size = 200,

```

```

        print_every = 100, verbose = True)
solver.train()
# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
#####
#                                     END OF YOUR CODE                                     #
#####

(Iteration 1 / 1225) loss: 2.301725
(Epoch 0 / 5) train acc: 0.187000; val_acc: 0.181000
(Iteration 101 / 1225) loss: 1.840451
(Iteration 201 / 1225) loss: 1.744794
(Epoch 1 / 5) train acc: 0.383000; val_acc: 0.418000
(Iteration 301 / 1225) loss: 1.529679
(Iteration 401 / 1225) loss: 1.543941
(Epoch 2 / 5) train acc: 0.444000; val_acc: 0.439000
(Iteration 501 / 1225) loss: 1.552313
(Iteration 601 / 1225) loss: 1.511572
(Iteration 701 / 1225) loss: 1.462240
(Epoch 3 / 5) train acc: 0.476000; val_acc: 0.444000
(Iteration 801 / 1225) loss: 1.515084
(Iteration 901 / 1225) loss: 1.462425
(Epoch 4 / 5) train acc: 0.495000; val_acc: 0.469000
(Iteration 1001 / 1225) loss: 1.376394
(Iteration 1101 / 1225) loss: 1.612014
(Iteration 1201 / 1225) loss: 1.612220
(Epoch 5 / 5) train acc: 0.509000; val_acc: 0.477000

```

10 Debug the training

With the default parameters we provided above, you should get a validation accuracy of about 0.36 on the validation set. This isn't very good.

One strategy for getting insight into what's wrong is to plot the loss function and the accuracies on the training and validation sets during optimization.

Another strategy is to visualize the weights that were learned in the first layer of the network. In most neural networks trained on visual data, the first layer weights typically show some visible structure when visualized.

[12]: *# Run this cell to visualize training loss and train / val accuracy*

```

plt.subplot(2, 1, 1)
plt.title('Training loss')
plt.plot(solver.loss_history, 'o')
plt.xlabel('Iteration')


plt.subplot(2, 1, 2)

```

```

plt.title('Accuracy')
plt.plot(solver.train_acc_history, '-o', label='train')
plt.plot(solver.val_acc_history, '-o', label='val')
plt.plot([0.5] * len(solver.val_acc_history), 'k--')
plt.xlabel('Epoch')
plt.legend(loc='lower right')
plt.gcf().set_size_inches(15, 12)
plt.show()

```



```


[13]: from cs231n.vis_utils import visualize_grid

# Visualize the weights of the network

def show_net_weights(net):
    W1 = net.params['W1']
    W1 = W1.reshape(3, 32, 32, -1).transpose(3, 1, 2, 0)
    plt.imshow(visualize_grid(W1, padding=3).astype('uint8'))
    plt.gca().axis('off')
    plt.show()

```

```
show_net_weights(model)
```



11 Tune your hyperparameters

What's wrong?. Looking at the visualizations above, we see that the loss is decreasing more or less linearly, which seems to suggest that the learning rate may be too low. Moreover, there is no gap between the training and validation accuracy, suggesting that the model we used has low capacity, and that we should increase its size. On the other hand, with a very large model we would expect to see more overfitting, which would manifest itself as a very large gap between the training and validation accuracy.

Tuning. Tuning the hyperparameters and developing intuition for how they affect the final performance is a large part of using Neural Networks, so we want you to get a lot of practice. Below, you should experiment with different values of the various hyperparameters, including hidden layer size, learning rate, number of training epochs, and regularization strength. You might also consider tuning the learning rate decay, but you should be able to get good performance using the default value.

Approximate results. You should be aim to achieve a classification accuracy of greater than 48% on the validation set. Our best network gets over 52% on the validation set.

Experiment: Your goal in this exercise is to get as good of a result on CIFAR-10 as you can (52% could serve as a reference), with a fully-connected Neural Network. Feel free implement your own techniques (e.g. PCA to reduce dimensionality, or adding dropout, or adding features to the solver, etc.).

```
[14]: best_model = None
      best_val = -1
      best_stats = []

#####
# TODO: Tune hyperparameters using the validation set. Store your best trained
#
# model in best_model.
#
#
# To help debug your network, it may help to use visualizations similar to the
# ones we used above; these visualizations will have significant qualitative
# differences from the ones we saw above for the poorly tuned network.
#
# Tweaking hyperparameters by hand can be fun, but you might find it useful to
# write code to sweep through possible combinations of hyperparameters
# automatically like we did on thes previous exercises.
#
#####
# *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
#https://github.com/MahanFathi/CS231/blob/master/assignment1/two_layer_net.
# ipynb

input_size = 32 * 32 * 3
num_classes = 10
```

```

lr_list = [1e-3, 1.5e-3]
reg_list = [2e0, 1.5e1]
hidden_size_list = [150, 180]
batch_size_list = [200, 256]
lr_decay_list = [0.97, 0.98]

for lr in lr_list:
    for reg in reg_list:
        for hidden_size in hidden_size_list:
            for batch_size in batch_size_list:
                for lr_decay in lr_decay_list:
                    # Create a two-layer network
                    model = TwoLayerNet(input_size, hidden_size, num_classes)
                    model.reg = reg

                    solver = Solver(model, data,
                                    update_rule = 'sgd_momentum',
                                    optim_config = {
                                        "learning_rate": lr,
                                    },
                                    lr_decay = lr_decay,
                                    num_epochs = 5, batch_size = batch_size,
                                    print_every = 100,
                                    verbose= False
                                    )

                    solver.train()
                    val_acc = solver.check_accuracy(data["X_val"],
                                                    data["y_val"],
                                                    num_samples = None,
                                                    batch_size = 200,
                                                    )

                    if val_acc > best_val:
                        best_val = val_acc
                        best_model = model
                        best_stats = solver

                    print('lr %e, reg %e, hid %d, batch_size %d, lr_decay %f, val_
↪accuracy: %f' % (
                        lr, reg, hidden_size, batch_size, lr_decay, val_acc))
                    #train accuracy: %f
print('best validation accuracy achieved: %f' % best_val)

# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
#####

```

```
#                                END OF YOUR CODE                                #
#####
```

```
lr 1.000000e-03, reg 2.000000e+00, hid 150, batch_size 200, lr_decay 0.970000,
val accuracy: 0.482000
lr 1.000000e-03, reg 2.000000e+00, hid 150, batch_size 200, lr_decay 0.980000,
val accuracy: 0.479000
lr 1.000000e-03, reg 2.000000e+00, hid 150, batch_size 256, lr_decay 0.970000,
val accuracy: 0.472000
lr 1.000000e-03, reg 2.000000e+00, hid 150, batch_size 256, lr_decay 0.980000,
val accuracy: 0.478000
lr 1.000000e-03, reg 2.000000e+00, hid 180, batch_size 200, lr_decay 0.970000,
val accuracy: 0.482000
lr 1.000000e-03, reg 2.000000e+00, hid 180, batch_size 200, lr_decay 0.980000,
val accuracy: 0.473000
lr 1.000000e-03, reg 2.000000e+00, hid 180, batch_size 256, lr_decay 0.970000,
val accuracy: 0.486000
lr 1.000000e-03, reg 2.000000e+00, hid 180, batch_size 256, lr_decay 0.980000,
val accuracy: 0.470000
lr 1.000000e-03, reg 1.500000e+01, hid 150, batch_size 200, lr_decay 0.970000,
val accuracy: 0.383000
lr 1.000000e-03, reg 1.500000e+01, hid 150, batch_size 200, lr_decay 0.980000,
val accuracy: 0.373000
lr 1.000000e-03, reg 1.500000e+01, hid 150, batch_size 256, lr_decay 0.970000,
val accuracy: 0.365000
lr 1.000000e-03, reg 1.500000e+01, hid 150, batch_size 256, lr_decay 0.980000,
val accuracy: 0.378000
lr 1.000000e-03, reg 1.500000e+01, hid 180, batch_size 200, lr_decay 0.970000,
val accuracy: 0.369000
lr 1.000000e-03, reg 1.500000e+01, hid 180, batch_size 200, lr_decay 0.980000,
val accuracy: 0.379000
lr 1.000000e-03, reg 1.500000e+01, hid 180, batch_size 256, lr_decay 0.970000,
val accuracy: 0.385000
lr 1.000000e-03, reg 1.500000e+01, hid 180, batch_size 256, lr_decay 0.980000,
val accuracy: 0.371000
lr 1.500000e-03, reg 2.000000e+00, hid 150, batch_size 200, lr_decay 0.970000,
val accuracy: 0.467000
lr 1.500000e-03, reg 2.000000e+00, hid 150, batch_size 200, lr_decay 0.980000,
val accuracy: 0.474000
lr 1.500000e-03, reg 2.000000e+00, hid 150, batch_size 256, lr_decay 0.970000,
val accuracy: 0.468000
lr 1.500000e-03, reg 2.000000e+00, hid 150, batch_size 256, lr_decay 0.980000,
val accuracy: 0.478000
lr 1.500000e-03, reg 2.000000e+00, hid 180, batch_size 200, lr_decay 0.970000,
val accuracy: 0.471000
lr 1.500000e-03, reg 2.000000e+00, hid 180, batch_size 200, lr_decay 0.980000,
val accuracy: 0.482000
```

```

lr 1.500000e-03, reg 2.000000e+00, hid 180, batch_size 256, lr_decay 0.970000,
val accuracy: 0.473000
lr 1.500000e-03, reg 2.000000e+00, hid 180, batch_size 256, lr_decay 0.980000,
val accuracy: 0.500000
lr 1.500000e-03, reg 1.500000e+01, hid 150, batch_size 200, lr_decay 0.970000,
val accuracy: 0.377000
lr 1.500000e-03, reg 1.500000e+01, hid 150, batch_size 200, lr_decay 0.980000,
val accuracy: 0.359000
lr 1.500000e-03, reg 1.500000e+01, hid 150, batch_size 256, lr_decay 0.970000,
val accuracy: 0.371000
lr 1.500000e-03, reg 1.500000e+01, hid 150, batch_size 256, lr_decay 0.980000,
val accuracy: 0.370000
lr 1.500000e-03, reg 1.500000e+01, hid 180, batch_size 200, lr_decay 0.970000,
val accuracy: 0.370000
lr 1.500000e-03, reg 1.500000e+01, hid 180, batch_size 200, lr_decay 0.980000,
val accuracy: 0.378000
lr 1.500000e-03, reg 1.500000e+01, hid 180, batch_size 256, lr_decay 0.970000,
val accuracy: 0.363000
lr 1.500000e-03, reg 1.500000e+01, hid 180, batch_size 256, lr_decay 0.980000,
val accuracy: 0.375000
best validation accuracy achieved: 0.500000

```

12 Test your model!

Run your best model on the validation and test sets. You should achieve above 48% accuracy on the validation set and the test set.

```

[15]: y_val_pred = np.argmax(best_model.loss(data['X_val']), axis=1)
      print('Validation set accuracy: ', (y_val_pred == data['y_val']).mean())

```

Validation set accuracy: 0.5

```

[16]: y_test_pred = np.argmax(best_model.loss(data['X_test']), axis=1)
      print('Test set accuracy: ', (y_test_pred == data['y_test']).mean())

```

Test set accuracy: 0.491

```


[17]: from cs231n.vis_utils import visualize_grid

      # Visualize the weights of the network

      def show_net_weights(net):
          W1 = net.params['W1']
          W1 = W1.reshape(3, 32, 32, -1).transpose(3, 1, 2, 0)
          plt.imshow(visualize_grid(W1, padding=3).astype('uint8'))
          plt.gca().axis('off')
          plt.show()

```

```
show_net_weights(model)
```



12.1 Inline Question 2:

Now that you have trained a Neural Network classifier, you may find that your testing accuracy is much lower than the training accuracy. In what ways can we decrease this gap? Select all that apply.

1. Train on a larger dataset.
2. Add more hidden units.
3. Increase the regularization strength.
4. None of the above.

Your Answer : We can decrease the gap between validation accuracy and testing accuracy by implementing: 1, 3 \ *Your Explanation :* Reasons are as following: 1. By training on the larger dataset, we are trying to incorporate greater variety of dataset. Increasing the dataset in a way that ensures the training and testing dataset comes from the same distribution will help improve testing accuracy. Hence, option 1 is the correct possible answer as long as train and test data has similar distribution. \

2. NOT a correct option. \

3. By increase the regularization strength, learned parameters are penalized against overfitting the data and hence improves the generalizability of the model which helps reduce the difference between training and testing accuracy. \

features

April 16, 2022

```
[ ]: # This mounts your Google Drive to the Colab VM.
from google.colab import drive
drive.mount('/content/drive')

# TODO: Enter the foldername in your Drive where you have saved the unzipped
# assignment folder, e.g. 'cs231n/assignments/assignment1/'
FOLDERNAME = 'Colab_Notebooks/cs231n/assignments/assignment1/'
assert FOLDERNAME is not None, "[!] Enter the foldername."

# Now that we've mounted your Drive, this ensures that
# the Python interpreter of the Colab VM can load
# python files from within it.
import sys
sys.path.append('/content/drive/My Drive/{}'.format(FOLDERNAME))

# This downloads the CIFAR-10 dataset to your Drive
# if it doesn't already exist.
%cd /content/drive/My\ Drive/$FOLDERNAME/cs231n/datasets/
!bash get_datasets.sh
%cd /content/drive/My\ Drive/$FOLDERNAME
```

```
Mounted at /content/drive
/content/drive/My
Drive/Colab_Notebooks/cs231n/assignments/assignment1/cs231n/datasets
/content/drive/My Drive/Colab_Notebooks/cs231n/assignments/assignment1
```

1 Image features exercise

Complete and hand in this completed worksheet (including its outputs and any supporting code outside of the worksheet) with your assignment submission. For more details see the [assignments page](#) on the course website.

We have seen that we can achieve reasonable performance on an image classification task by training a linear classifier on the pixels of the input image. In this exercise we will show that we can improve our classification performance by training linear classifiers not on raw pixels but on features that are computed from the raw pixels.

All of your work for this exercise will be done in this notebook.

```
[ ]: import random
import numpy as np
from cs231n.data_utils import load_CIFAR10
import matplotlib.pyplot as plt

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/
# → autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2
```

1.1 Load data

Similar to previous exercises, we will load CIFAR-10 data from disk.

```
[ ]: from cs231n.features import color_histogram_hsv, hog_feature

def get_CIFAR10_data(num_training=49000, num_validation=1000, num_test=1000):
    # Load the raw CIFAR-10 data
    cifar10_dir = 'cs231n/datasets/cifar-10-batches-py'

    # Cleaning up variables to prevent loading data multiple times (which may
    # → cause memory issue)
    try:
        del X_train, y_train
        del X_test, y_test
        print('Clear previously loaded data.')
    except:
        pass

    X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)

    # Subsample the data
    mask = list(range(num_training, num_training + num_validation))
    X_val = X_train[mask]
    y_val = y_train[mask]
    mask = list(range(num_training))
    X_train = X_train[mask]
```



```

y_train = y_train[mask]
mask = list(range(num_test))
X_test = X_test[mask]
y_test = y_test[mask]

return X_train, y_train, X_val, y_val, X_test, y_test

X_train, y_train, X_val, y_val, X_test, y_test = get_CIFAR10_data()

```

1.2 Extract Features

For each image we will compute a Histogram of Oriented Gradients (HOG) as well as a color histogram using the hue channel in HSV color space. We form our final feature vector for each image by concatenating the HOG and color histogram feature vectors.

Roughly speaking, HOG should capture the texture of the image while ignoring color information, and the color histogram represents the color of the input image while ignoring texture. As a result, we expect that using both together ought to work better than using either alone. Verifying this assumption would be a good thing to try for your own interest.

The `hog_feature` and `color_histogram_hsv` functions both operate on a single image and return a feature vector for that image. The `extract_features` function takes a set of images and a list of feature functions and evaluates each feature function on each image, storing the results in a matrix where each column is the concatenation of all feature vectors for a single image.

```

[ ]: from cs231n.features import *

num_color_bins = 10 # Number of bins in the color histogram
feature_fns = [hog_feature, lambda img: color_histogram_hsv(img,
    ↪nbin=num_color_bins)]
X_train_feats = extract_features(X_train, feature_fns, verbose=True)
X_val_feats = extract_features(X_val, feature_fns)
X_test_feats = extract_features(X_test, feature_fns)

# Preprocessing: Subtract the mean feature
mean_feat = np.mean(X_train_feats, axis=0, keepdims=True)
X_train_feats -= mean_feat
X_val_feats -= mean_feat
X_test_feats -= mean_feat

# Preprocessing: Divide by standard deviation. This ensures that each feature
# has roughly the same scale.
std_feat = np.std(X_train_feats, axis=0, keepdims=True)
X_train_feats /= std_feat
X_val_feats /= std_feat
X_test_feats /= std_feat

```

```
# Preprocessing: Add a bias dimension
```

```
X_train_feats = np.hstack([X_train_feats, np.ones((X_train_feats.shape[0], 1))])
```

```
X_val_feats = np.hstack([X_val_feats, np.ones((X_val_feats.shape[0], 1))])
```

```
X_test_feats = np.hstack([X_test_feats, np.ones((X_test_feats.shape[0], 1))])
```

```
Done extracting features for 1000 / 49000 images
Done extracting features for 2000 / 49000 images
Done extracting features for 3000 / 49000 images
Done extracting features for 4000 / 49000 images
Done extracting features for 5000 / 49000 images
Done extracting features for 6000 / 49000 images
Done extracting features for 7000 / 49000 images
Done extracting features for 8000 / 49000 images
Done extracting features for 9000 / 49000 images
Done extracting features for 10000 / 49000 images
Done extracting features for 11000 / 49000 images
Done extracting features for 12000 / 49000 images
Done extracting features for 13000 / 49000 images
Done extracting features for 14000 / 49000 images
Done extracting features for 15000 / 49000 images
Done extracting features for 16000 / 49000 images
Done extracting features for 17000 / 49000 images
Done extracting features for 18000 / 49000 images
Done extracting features for 19000 / 49000 images
Done extracting features for 20000 / 49000 images
Done extracting features for 21000 / 49000 images
Done extracting features for 22000 / 49000 images
Done extracting features for 23000 / 49000 images
Done extracting features for 24000 / 49000 images
Done extracting features for 25000 / 49000 images
Done extracting features for 26000 / 49000 images
Done extracting features for 27000 / 49000 images
Done extracting features for 28000 / 49000 images
Done extracting features for 29000 / 49000 images
Done extracting features for 30000 / 49000 images
Done extracting features for 31000 / 49000 images
Done extracting features for 32000 / 49000 images
Done extracting features for 33000 / 49000 images
Done extracting features for 34000 / 49000 images
Done extracting features for 35000 / 49000 images
Done extracting features for 36000 / 49000 images
Done extracting features for 37000 / 49000 images
Done extracting features for 38000 / 49000 images
Done extracting features for 39000 / 49000 images
Done extracting features for 40000 / 49000 images
Done extracting features for 41000 / 49000 images
Done extracting features for 42000 / 49000 images
```

```

Done extracting features for 43000 / 49000 images
Done extracting features for 44000 / 49000 images
Done extracting features for 45000 / 49000 images
Done extracting features for 46000 / 49000 images
Done extracting features for 47000 / 49000 images
Done extracting features for 48000 / 49000 images
Done extracting features for 49000 / 49000 images

```

1.3 Train SVM on features

Using the multiclass SVM code developed earlier in the assignment, train SVMs on top of the features extracted above; this should achieve better results than training SVMs directly on top of raw pixels.

```

[ ]: # Use the validation set to tune the learning rate and regularization strength

from cs231n.classifiers.linear_classifier import LinearSVM

learning_rates = [1e-7, 5e-7, 8e-7]
regularization_strengths = [5e4, 7e4, 5e5]

results = {}
best_val = -1
best_svm = None

#####
# TODO:
# Use the validation set to set the learning rate and regularization strength. #
# This should be identical to the validation that you did for the SVM; save #
# the best trained classifier in best_svm. You might also want to play #
# with different numbers of bins in the color histogram. If you are careful #
# you should be able to get accuracy of near 0.44 on the validation set. #
#####
# *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

for lr in learning_rates:
    for reg in regularization_strengths:
        curr_svm = LinearSVM()
        loss_hist = curr_svm.train(X_train_feats, y_train,
                                   learning_rate=lr, reg=reg,
                                   num_iters=1500, verbose=False)

        y_train_pred = curr_svm.predict(X_train_feats)
        train_accuracy = np.mean(y_train == y_train_pred)

        y_val_pred = curr_svm.predict(X_val_feats)
        val_accuracy = np.mean(y_val == y_val_pred)

```

```

if (val_accuracy > best_val):
    best_val = val_accuracy
    best_svm = curr_svm

results[(lr, reg)] = (train_accuracy, val_accuracy)

# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

# Print out results.
for lr, reg in sorted(results):
    train_accuracy, val_accuracy = results[(lr, reg)]
    print('lr %e reg %e train accuracy: %f val accuracy: %f' % (
        lr, reg, train_accuracy, val_accuracy))

print('best validation accuracy achieved: %f' % best_val)

```

```

lr 1.000000e-07 reg 5.000000e+04 train accuracy: 0.417327 val accuracy: 0.425000
lr 1.000000e-07 reg 7.000000e+04 train accuracy: 0.418347 val accuracy: 0.419000
lr 1.000000e-07 reg 5.000000e+05 train accuracy: 0.412837 val accuracy: 0.422000
lr 5.000000e-07 reg 5.000000e+04 train accuracy: 0.415776 val accuracy: 0.411000
lr 5.000000e-07 reg 7.000000e+04 train accuracy: 0.411102 val accuracy: 0.401000
lr 5.000000e-07 reg 5.000000e+05 train accuracy: 0.383469 val accuracy: 0.385000
lr 8.000000e-07 reg 5.000000e+04 train accuracy: 0.409776 val accuracy: 0.409000
lr 8.000000e-07 reg 7.000000e+04 train accuracy: 0.418898 val accuracy: 0.419000
lr 8.000000e-07 reg 5.000000e+05 train accuracy: 0.369857 val accuracy: 0.392000
best validation accuracy achieved: 0.425000

```

```

[ ]: # Evaluate your trained SVM on the test set: you should be able to get at least
    ↪ 0.40
y_test_pred = best_svm.predict(X_test_feats)
test_accuracy = np.mean(y_test == y_test_pred)
print(test_accuracy)

```

0.418

```

[ ]: # An important way to gain intuition about how an algorithm works is to
# visualize the mistakes that it makes. In this visualization, we show examples
# of images that are misclassified by our current system. The first column
# shows images that our system labeled as "plane" but whose true label is
# something other than "plane".

examples_per_class = 8
classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse',
    ↪ 'ship', 'truck']

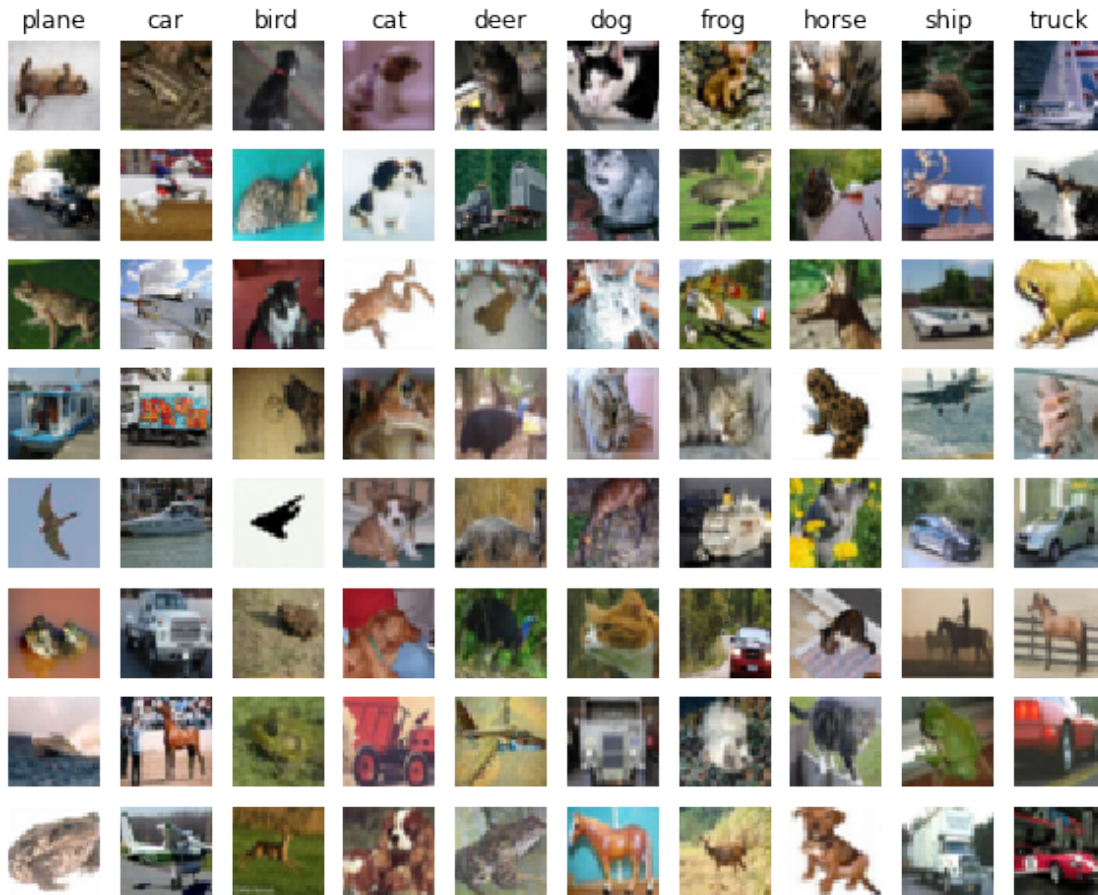
for cls, cls_name in enumerate(classes):
    idxs = np.where((y_test != cls) & (y_test_pred == cls))[0]

```

```

idxs = np.random.choice(idxs, examples_per_class, replace=False)
for i, idx in enumerate(idxs):
    plt.subplot(examples_per_class, len(classes), i * len(classes) + cls + 1)
    plt.imshow(X_test[idx].astype('uint8'))
    plt.axis('off')
    if i == 0:
        plt.title(cls_name)
plt.show()

```



1.3.1 Inline question 1:

Describe the misclassification results that you see. Do they make sense?

Your Answer : \ Those are mainly features in high dimensional space. For example: one of the horse is classified as dog and on the other hand one dog is classified as horse given both of the animals have 4 legs, tail etc. So it would be important to extract and analyzing the intermediate layers from the high dimensional ones.

1.4 Neural Network on image features

Earlier in this assignment we saw that training a two-layer neural network on raw pixels achieved better classification performance than linear classifiers on raw pixels. In this notebook we have seen that linear classifiers on image features outperform linear classifiers on raw pixels.

For completeness, we should also try training a neural network on image features. This approach should outperform all previous approaches: you should easily be able to achieve over 55% classification accuracy on the test set; our best model achieves about 60% classification accuracy.

```
[ ]: # Preprocessing: Remove the bias dimension
# Make sure to run this cell only ONCE
print(X_train_feats.shape)
X_train_feats = X_train_feats[:, :-1]
X_val_feats = X_val_feats[:, :-1]
X_test_feats = X_test_feats[:, :-1]

print(X_train_feats.shape)
```

```
(49000, 155)
```

```
(49000, 154)
```

```
[ ]: from cs231n.classifiers.fc_net import TwoLayerNet
from cs231n.solver import Solver

input_dim = X_train_feats.shape[1]
hidden_dim = 500
num_classes = 10

data = {
    'X_train': X_train_feats,
    'y_train': y_train,
    'X_val': X_val_feats,
    'y_val': y_val,
    'X_test': X_test_feats,
    'y_test': y_test,
}

net = TwoLayerNet(input_dim, hidden_dim, num_classes)
best_net = None

#####
# TODO: Train a two-layer neural network on image features. You may want to #
# cross-validate various parameters as in previous sections. Store your best #
# model in the best_net variable.                                           #
#####
# *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
```

```

lr_list = [5e-1, 7e-1, 8e-1]
reg_list = [2e-1, 1.5e0]
hidden_size_list = [500, 150, 180]
batch_size_list = [200, 256]
lr_decay_list = [0.97, 0.98]

for lr in lr_list:
    for reg in reg_list:
        #for hidden_dim in hidden_size_list:
        for batch_size in batch_size_list:
            for lr_decay in lr_decay_list:
                # Create a two-layer network
                net = TwoLayerNet(input_dim, hidden_dim, num_classes)
                net.reg = reg

                solver = Solver(net, data,
                                update_rule = 'sgd_momentum',
                                optim_config = {
                                    "learning_rate": lr,
                                },
                                lr_decay = lr_decay,
                                num_epochs = 5, batch_size = batch_size,
                                print_every = 100,
                                verbose= False
                                )

                solver.train()
                val_acc = solver.check_accuracy(data["X_val"],
                                                data["y_val"],
                                                num_samples = None,
                                                batch_size = 200,
                                                )

                if val_acc > best_val:
                    best_val = val_acc
                    best_net = net
                    best_stats = solver
                print('lr %e, reg %e, hid %d, batch_size %d, lr_decay %f, val_
↪accuracy: %f' % (
                    lr, reg, hidden_dim, batch_size, lr_decay, val_acc))
                #train accuracy: %f
print('best validation accuracy achieved: %f' % best_val)

# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

```

lr 5.000000e-01, reg 2.000000e-01, hid 500, batch_size 200, lr_decay 0.970000,
val accuracy: 0.587000

lr 5.000000e-01, reg 2.000000e-01, hid 500, batch_size 200, lr_decay 0.980000,

```
val accuracy: 0.585000
lr 5.000000e-01, reg 2.000000e-01, hid 500, batch_size 256, lr_decay 0.970000,
val accuracy: 0.594000
lr 5.000000e-01, reg 2.000000e-01, hid 500, batch_size 256, lr_decay 0.980000,
val accuracy: 0.568000
lr 5.000000e-01, reg 1.500000e+00, hid 500, batch_size 200, lr_decay 0.970000,
val accuracy: 0.578000
lr 5.000000e-01, reg 1.500000e+00, hid 500, batch_size 200, lr_decay 0.980000,
val accuracy: 0.581000
lr 5.000000e-01, reg 1.500000e+00, hid 500, batch_size 256, lr_decay 0.970000,
val accuracy: 0.588000
lr 5.000000e-01, reg 1.500000e+00, hid 500, batch_size 256, lr_decay 0.980000,
val accuracy: 0.584000
lr 7.000000e-01, reg 2.000000e-01, hid 500, batch_size 200, lr_decay 0.970000,
val accuracy: 0.579000
lr 7.000000e-01, reg 2.000000e-01, hid 500, batch_size 200, lr_decay 0.980000,
val accuracy: 0.576000
lr 7.000000e-01, reg 2.000000e-01, hid 500, batch_size 256, lr_decay 0.970000,
val accuracy: 0.573000
lr 7.000000e-01, reg 2.000000e-01, hid 500, batch_size 256, lr_decay 0.980000,
val accuracy: 0.585000
lr 7.000000e-01, reg 1.500000e+00, hid 500, batch_size 200, lr_decay 0.970000,
val accuracy: 0.568000
lr 7.000000e-01, reg 1.500000e+00, hid 500, batch_size 200, lr_decay 0.980000,
val accuracy: 0.578000
lr 7.000000e-01, reg 1.500000e+00, hid 500, batch_size 256, lr_decay 0.970000,
val accuracy: 0.574000
lr 7.000000e-01, reg 1.500000e+00, hid 500, batch_size 256, lr_decay 0.980000,
val accuracy: 0.585000
lr 8.000000e-01, reg 2.000000e-01, hid 500, batch_size 200, lr_decay 0.970000,
val accuracy: 0.589000
lr 8.000000e-01, reg 2.000000e-01, hid 500, batch_size 200, lr_decay 0.980000,
val accuracy: 0.588000
lr 8.000000e-01, reg 2.000000e-01, hid 500, batch_size 256, lr_decay 0.970000,
val accuracy: 0.566000
lr 8.000000e-01, reg 2.000000e-01, hid 500, batch_size 256, lr_decay 0.980000,
val accuracy: 0.565000
lr 8.000000e-01, reg 1.500000e+00, hid 500, batch_size 200, lr_decay 0.970000,
val accuracy: 0.574000
lr 8.000000e-01, reg 1.500000e+00, hid 500, batch_size 200, lr_decay 0.980000,
val accuracy: 0.580000
lr 8.000000e-01, reg 1.500000e+00, hid 500, batch_size 256, lr_decay 0.970000,
val accuracy: 0.585000
lr 8.000000e-01, reg 1.500000e+00, hid 500, batch_size 256, lr_decay 0.980000,
val accuracy: 0.577000
best validation accuracy achieved: 0.594000
```



```
[ ]: # Run your best neural net classifier on the test set. You should be able  
# to get more than 55% accuracy.
```

```
y_test_pred = np.argmax(best_net.loss(data['X_test']), axis=1)  
test_acc = (y_test_pred == data['y_test']).mean()  
print(test_acc)
```

0.565